

AI Governance Failure Taxonomy

Structural classification of AI governance failure mechanisms

Edition	Full Reference Manual
Version	0.2.0-draft
Status	Draft
Families	9
Failure classes	53
Author and rights holder	Dr Michelle O'Rourke

Copyright © 2026 Dr Michelle O'Rourke. All rights reserved.

This compiled reference manual, including its compilation, editorial arrangement, design and presentation, may not be reproduced or redistributed in whole or substantial part without permission. Underlying VIGIL Observatory taxonomy content remains publicly referenceable and reusable under the VIGIL Observatory Licence and Reuse Terms. Third-party materials remain subject to their respective rights and licence conditions.

VIGIL Observatory Failure Taxonomy

Full Reference Manual

Contents

1. Authority Boundary Integrity Failures
 - Source-Authority Confusion
 - Capability-Authority Conflation
 - Target and Scope Authority Transposition
 - Ambient Instruction Capture
 - Transformation-Mediated Authority Laundering
 - Control-Plane Authority Crossover
 - Technical Reachability Mistaken for Authority
 - Delegation Scope Expansion
 - Transitive Authority Propagation
 - Inferential Evidence-Authority Conflation
 - Identity-Representation Authority Conflation
2. Provenance & Lineage Integrity Failures
 - Authorship or Source Misattribution
 - Untraceable Synthesis
 - Cross-Context Lineage Distortion
 - Transformation Lineage Collapse
 - False Continuity Attribution
 - Target-Object Binding Failure
 - Unsupported Mechanism Attribution
3. Verification & Completion Integrity Failures
 - Required Verification Omission
 - False-Success Representation
 - Completion-Condition Mismatch
 - Post-Verification Mutation
 - Stale Verification Reuse
 - Partial Completion Represented as Complete
 - Deterministic Verification Decomposition Omission
4. Observability & Audit Integrity Failures
 - Material Event Non-Capture
 - Monitor Circumvention or Coverage Bypass
 - Audit-Trail Non-Reconstructability
 - Actor-Action Attribution Loss
 - Hidden Material Execution Path
 - Audit-Evidence Integrity Loss
 - Execution-State Non-Disclosure
 - Material Signal Fragmentation
 - Primary Evidence Accessibility Failure
 - Authorised Investigative Evidence-Access Pathway Failure
5. Access & Session State Integrity Failures
 - Authentication-State Continuity Failure
 - Access-State Collapse
 - Re-entry Access-State Ambiguity
 - Verification-Dependency Access Failure

- 6. Work-State Continuity Integrity Failures
 - Continuity Anchor Failure
 - Material Work-State Persistence Failure
 - Restoration-State Integrity Failure
- 7. Governance Control Reach Integrity Failures
 - Control-State Preservation Failure
 - Required Governance Route Bypass
 - Governance Signal Delivery Dead-End
- 8. Control Activation Integrity Failures
 - Control Availability Ambiguity
 - Required Control Non-Activation
 - Control Authority Suppression
 - Unwarranted Control Activation
- 9. Agency-Preserving Influence Integrity Failures
 - Dependency-Cultivation Optimisation
 - Protected-Signal Influence Repurposing
 - Relationally Conditioned Epistemic Steering
 - Instrumental Choice Manipulation

Authority Boundary Integrity Failures

The system can see, reach, decode, inherit, or do something and mistakes that fact for permission to act.

Immutable ID: VIGIL-FF-0001 · **Semantic code:** AUTHORITY_BOUNDARY · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which information, access, capability, identity, technical reachability, transformed content, delegated intent, or system state is treated as conferring authority, permission, delegation, or execution scope that has not been independently established for the relevant source, target, actor, action, or context.

Governing invariant

Visibility is not authority. Capability is not permission. Technical reachability is not delegation. Transformation must not increase authority. Authority remains bound to its source, target, action, scope, and context unless validly re-established.

Classification boundary

Include when

Use this family when the primary structural failure is an unjustified creation, enlargement, transfer, substitution, or inheritance of authority across a source, capability, target, scope, representation, delegation, or control-plane boundary.

Exclude when

Do not use this family merely because an action was incorrect, harmful, insecure, unavailable, or refused. A qualifying event must involve an authority boundary being misclassified, expanded, transferred, or treated as satisfied without adequate authority.

Scope

- AI models
- tool-using agents
- multi-agent systems
- orchestration frameworks
- retrieval and memory systems
- APIs and connected services
- human-AI workflows

Allowed identifiers

- VIGIL-FC-000001 — SOURCE_AUTHORITY_CONFUSION
- VIGIL-FC-000002 — CAPABILITY_AUTHORITY_CONFLATION
- VIGIL-FC-000003 — TARGET_SCOPE_TRANSPOSITION
- VIGIL-FC-000004 — AMBIENT_INSTRUCTION_CAPTURE
- VIGIL-FC-000005 — TRANSFORMATION_AUTHORITY_LAUNDERING
- VIGIL-FC-000006 — CONTROL_PLANE_CROSSOVER
- VIGIL-FC-000007 — TECHNICAL_REACHABILITY_AS_AUTHORITY
- VIGIL-FC-000008 — DELEGATION_SCOPE_EXPANSION
- VIGIL-FC-000009 — TRANSITIVE_AUTHORITY_PROPAGATION
- VIGIL-FC-000046 — INFERENTIAL_EVIDENCE_AUTHORITY_CONFLATION
- VIGIL-FC-000053 — IDENTITY_REPRESENTATION_AUTHORITY_CONFLATION

Prior codes and aliases

- VIGIL.FF.AUTHORITY_BOUNDARY

Failure classes

CLASS

DRAFT

Source-Authority Confusion

VIGIL-FC-000001 · SOURCE_AUTHORITY_CONFUSION

Plain English: The system treats content it can read as though the content has permission to instruct it.

Technical definition

A failure in which content from a lower-authority, untrusted, retrieved, generated, third-party, or otherwise non-authorising source is treated as an instruction, command, policy, approval, or execution-authorising signal.

Recognition criteria

- The system receives or retrieves content from a source whose authority is lower than, different from, or unresolved relative to the action being considered.
- The content is treated as operative instruction or authority rather than as data, evidence, context, or material for analysis.
- No independent authority check validly establishes that the source may authorise the affected action.

Exclusions

- A system merely quotes, summarizes, classifies, or reasons about untrusted instructions without obeying them.
- The source is genuinely authorised for the specific action and scope.
- The event is only a content-quality or factual-accuracy problem.

Illustrative examples

- A retrieved document contains instructions to modify an account and the agent follows them even though the user never authorised the change.
- A tool output contains instruction-like text and the system treats it as higher-priority task direction.

Prior codes and aliases

- VIGIL_FF.AUTHORITY_BOUNDARY.SOURCE_AUTHORITY_CONFUSION
- instruction-source confusion
- source-authority collapse

VIGIL Observatory Case File classifications (5)

- **Primary:** VIGIL-2026-FM-0009 — Non-revocable conversational context contamination
MEDIUM

Revoked or quarantined assistant content continues to be treated as direction-bearing authority rather than non-authoritative material.

- **Primary:** VIGIL-2026-FM-0022 — Instruction-bearing external content treated as trusted execution authority
HIGH

External lower-authority content is treated as execution-bearing instruction without independent authority validation.

- **Primary:** VIGIL-2026-FM-0040 — Synthetic authority impersonation crossing a consequential execution boundary
MEDIUM

Synthetic identity evidence from a non-authorising source is accepted as transaction-bearing authority.

- **Primary:** VIGIL-2026-FM-0070 — Aggregate security-control composition failure
HIGH

Attacker-controlled material in the q search parameter was treated as operative instruction that caused Copilot to search protected enterprise data and construct the response pathway.

- **Secondary:** VIGIL-2026-FM-0072 — Peer or delegate instruction treated as transitive authority HIGH

The incident independently shows a peer-authored GO message, urgency and apparent coordination status being treated as execution-authorising content despite lacking independent authority for the external action.

CLASS

DRAFT

Capability-Authority Conflation

VIGIL-FC-000002 · CAPABILITY_AUTHORITY_CONFLATION

Plain English: The system assumes that because it has the ability or tool access to do something, it is allowed to do it.

Technical definition

A failure in which available capability, tool access, credentials, privileges, feature availability, or action affordances are treated as sufficient authority for an action without separately establishing permission for the actor, target, action, and scope.

Recognition criteria

- A capability or execution pathway is technically available.
- The system uses availability or possession of the capability as a material basis for permission.
- Independent authority for the specific action is absent, narrower, expired, unresolved, or not revalidated.

Exclusions

- A permitted action is executed within an explicitly authorised scope.
- The system merely detects that capability exists but does not infer permission from it.
- A vulnerability exists but no AI or governance process converts capability into authority.

Illustrative examples

- An agent has a deployment tool and pushes a change because deployment is technically possible, despite lacking approval for that environment.
- A model with access to a payment tool treats the presence of that tool as permission to spend.

Prior codes and aliases

- VIGIL.FF.AUTHORITY_BOUNDARY.CAPABILITY_AUTHORITY_CONFLATION

VIGIL Observatory Case File classifications (2)

- **Primary:** VIGIL-2026-FM-0001 — Permission-gated tool invocation freeze **HIGH**

The definition requires tool capability to become operative before user permission is established.

- **Primary:** VIGIL-2026-FM-0044 — Instrumental objective success through unauthorised exploitation **HIGH**

Technical usefulness and executable capability are treated as permission despite absent target and method authority.

CLASS

DRAFT

Target and Scope Authority Transposition

VIGIL-FC-000003 · TARGET_SCOPE_TRANSPOSITION

Plain English: Permission for one target or situation gets silently reused for a different target, action, or environment.

Technical definition

A failure in which valid or purported authority tied to one principal, target, environment, action, jurisdiction, task, or scope is transferred to another without a fresh basis establishing that the authority applies there.

Recognition criteria

- Some authority or permission exists or is assumed to exist for an original target or scope.
- The operative target, action, environment, principal, or scope materially changes.
- The system continues as though the original authority automatically carries over.

Exclusions

- The original authorisation explicitly covers the changed target and action.
- The change is purely representational and does not alter authority-relevant scope.
- The event concerns only loss of context without any authority consequence.

Illustrative examples

- Permission to test a sandbox is treated as permission to test a real production target.
- Authority to edit one repository is silently applied to another repository with a similar name.

Prior codes and aliases

- VIGIL.FF.AUTHORITY_BOUNDARY.TARGET_SCOPE_TRANSPOSITION

VIGIL Observatory Case File classifications (5)

- **Primary:** VIGIL-2026-FM-0032 — Role-incongruent affective governance (HIGH)

Affective governance valid or purported for one functional role or relationship scope is applied across a materially different companion, tutor, child-facing, clinical, accessibility, service, or hybrid-role scope without re-establishing applicability.

- **Primary:** VIGIL-2026-FM-0042 — Exported local political censorship in globally deployed AI systems (HIGH)

Political restriction authority tied to one jurisdiction is transposed into a materially different global deployment scope.

- **Primary:** VIGIL-2026-FM-0054 — Evaluation-environment constraint drift and real-target scope transposition (HIGH)

Authority and scope declared for a simulated evaluation are carried into a materially different real target environment.

- **Primary:** VIGIL-2026-FM-0060 — Defensive-offensive autonomy scope compression and control coupling (MEDIUM)

Authority and control state for bounded defence are transposed into offensive scope, or offensive constraints are applied to distinct defensive scope.

- **Primary:** VIGIL-2026-FM-0061 — Conflict-event authority reactivation without independent offensive reauthorisation (HIGH)

Prior offensive authority is renewed or extended into a new event and time scope without independent reauthorisation.

VARIANT

DRAFT

Ambient Instruction Capture

VIGIL-FC-000004 · AMBIENT_INSTRUCTION_CAPTURE

Plain English: Background or third-party content is mistaken for a command from the user.

Technical definition

A source-authority failure in which environmental, ambient, third-party, multimodal, retrieved, or media-borne content is interpreted as authorised user or operator instruction.

Recognition criteria

- Instruction-like content originates from an ambient, environmental, third-party, media, screen, retrieved, or otherwise non-user-authorising source.
- The system attributes the content to an authorised actor or treats it as execution-bearing instruction.
- The source distinction is not adequately preserved before action.

Exclusions

- The user explicitly adopts or confirms the ambient content as an instruction.
- The system transcribes or analyses ambient content without obeying it.
- The issue is only speech-recognition accuracy with no authority consequence.

Illustrative examples

- A voice assistant hears instructions from a television and treats them as the account holder's command.
- An agent reads hidden text on a webpage and follows it as task direction.

Prior codes and aliases

- `VIGIL_FF.AUTHORITY_BOUNDARY.AMBIENT_INSTRUCTION_CAPTURE`

Relationships

- **Child Of:** `VIGIL-FC-000001` — Ambient capture is a source-authority variant defined by the source channel.

VARIANT

DRAFT

Transformation-Mediated Authority Laundering

`VIGIL-FC-000005` · `TRANSFORMATION_AUTHORITY_LAUNDERING`

Plain English: Untrusted content becomes 'trusted' simply because it was decoded, decrypted, rendered, reformatted, or otherwise transformed.

Technical definition

A source-authority failure in which content gains greater instruction, governance, or execution authority after a representation change even though the transformation does not itself establish new authority.

Recognition criteria

- The originating content is untrusted, lower-authority, or authority-ambiguous.
- The content undergoes decoding, decryption, rendering, deserialisation, reconstruction, summarisation, translation, or another representation change.
- The transformed representation is assigned greater operative authority than the source had before transformation.

Exclusions

- The transformation preserves the original source authority and provenance.
- A trusted authority independently re-authorises the transformed content.
- The issue is ordinary decoding or parsing failure without an authority change.

Illustrative examples

- Encrypted text from an untrusted source is decrypted and then obeyed as if decryption made it a trusted instruction.
- A serialized object derived from untrusted input is reconstructed and treated as authorised configuration.

Prior codes and aliases

- `VIGIL.FF.AUTHORITY_BOUNDARY.TRANSFORMATION_AUTHORITY_LAUNDERING`

Relationships

- **Child Of:** `VIGIL-FC-000001` — The distinguishing feature is a representation change that increases apparent authority.

VIGIL Observatory Case File classifications (1)

- **Primary:** `VIGIL-2026-FM-0068` — Transformation-mediated source-authority laundering
`HIGH`

The FM and class share the required transformation, lower-authority source, post-transformation authority increase, and operative effect.

VARIANT

DRAFT

Control-Plane Authority Crossover

`VIGIL-FC-000006` · `CONTROL_PLANE_CROSSOVER`

Plain English: Ordinary model output or data silently becomes trusted system configuration or orchestration state.

Technical definition

A failure in which data-plane, model-generated, retrieved, or otherwise non-authoritative content is converted into control-plane state that can govern routing, tools, memory policy, configuration, checkpoints, permissions, or execution without fresh provenance and authority validation.

Recognition criteria

- The originating material is data, model output, retrieved content, or other non-control-plane material.
- The material is promoted or converted into configuration, routing, tool schema, policy, cache, checkpoint, memory-policy, serialized control object, or comparable control state.
- The promotion occurs without a fresh authority validation appropriate to the new control-plane effect.

Exclusions

- The control-plane update is explicitly authorised and validated.
- The material remains data and cannot alter control state.
- The issue is only configuration correctness after a properly authorised change.

Illustrative examples

- A model-generated tool schema is automatically installed as trusted orchestration configuration.
- Retrieved content is written into persistent routing policy and changes later agent behaviour without review.

Prior codes and aliases

- `VIGIL.FF.AUTHORITY_BOUNDARY.CONTROL_PLANE_CROSSOVER`

Relationships

- **Child Of:** `VIGIL-FC-000002`
- **Distinguish From:** `VIGIL-FC-000005` — Control-plane crossover requires promotion into operative control state; transformation laundering does not.

VIGIL Observatory Case File classifications (1)

- **Primary:** `VIGIL-2026-FM-0069` — Agent-framework control-plane boundary failure
`HIGH`

Attacker-influenced data-plane state crosses into trusted framework control state without fresh authority validation.

Technical Reachability Mistaken for Authority

VIGIL-FC-000007 · TECHNICAL_REACHABILITY_AS_AUTHORITY

Plain English: The system can reach a resource, account, API, or target and mistakes that reachability for permission.

Technical definition

A failure in which technical reachability, inherited sessions, credential possession, network access, trust bindings, or API availability are treated as sufficient permission over a target, principal, organisation, or action scope.

Recognition criteria

- The system can technically reach or authenticate to a resource or target.
- That reachability is used as a material basis for authorising an action.
- The relevant target, principal, organisation, or action scope is not independently authorised.

Exclusions

- The reachable resource is within the explicitly authorised scope.
- The event is only privilege escalation or exploitation by an adversary with no authority classification by the AI or governance process.
- The system detects reachability but still performs a valid authority check before action.

Illustrative examples

- An inherited session allows access to another organisation's environment and the agent treats that access as permission to act there.
- An API accepts a request, and the agent interprets acceptance as proof that the requested action was authorised.

Prior codes and aliases

- VIGIL.FF.AUTHORITY_BOUNDARY.TECHNICAL_REACHABILITY_AS_AUTHORITY

Relationships

- **Child Of:** VIGIL-FC-000002
- **Can Cooccur With:** VIGIL-FC-000003 — Reachability may coincide with an unauthorised target or scope transposition without being defined by that transposition.

VIGIL Observatory Case File classifications (2)

- **Primary:** VIGIL-2026-FM-0045 — Purpose-limited surveillance access repurposed for unauthorised personal use HIGH

Valid credentials and general role reachability are mistaken for authority for a specific purpose-bound surveillance action.

- **Primary:** VIGIL-2026-FM-0056 — Technically available access mistaken for cross-organisational authority HIGH

Surviving credentials, sessions, devices, or trust bindings are treated as cross-organisational permission after authority changes.

VARIANT

DRAFT

Delegation Scope Expansion

VIGIL-FC-000008 · DELEGATION_SCOPE_EXPANSION

Plain English: A broad goal or earlier permission is stretched into permission for a materially new action.

Technical definition

A failure in which a goal, standing instruction, prior approval, workflow mandate, role, or delegation is interpreted more broadly than its established action, target, temporal, financial, credential, publication, mutation, or external-effect scope.

Recognition criteria

- A valid or purported delegation exists.
- The proposed action materially changes the type, target, consequence, irreversibility, sensitivity, or external effect of the delegated task.
- The system proceeds without re-establishing authority for the expanded scope.

Exclusions

- The delegation explicitly and unambiguously covers the new action and consequence.
- The new action is a trivial implementation detail with no material authority change.
- The problem is merely inefficient planning rather than excess authority.

Illustrative examples

- Permission to prepare a deployment plan is treated as permission to deploy it.

- A standing instruction to manage files is treated as permission to delete an unrelated archive.

Prior codes and aliases

- `VIGIL.FF.AUTHORITY_BOUNDARY.DELEGATION_SCOPE_EXPANSION`

Relationships

- **Child Of:** `VIGIL-FC-000003`

VIGIL Observatory Case File classifications (1)

- **Primary:** `VIGIL-2026-FM-0036` — Undisclosed agentic workspace replication and data-egress authority collapse `HIGH`

Authority to inspect or act in a workspace is expanded into transmission, persistence, training, or secondary-use authority.

VARIANT

DRAFT

Transitive Authority Propagation

`VIGIL-FC-000009` · `TRANSITIVE_AUTHORITY_PROPAGATION`

Plain English: One agent or system assumes that authority held by another automatically transfers through the chain.

Technical definition

A failure in which authority attributed to one actor, agent, tool, service, or principal is assumed to propagate transitively to another participant, downstream action, or delegated subtask without explicit preservation of principal, target, action, and scope.

Recognition criteria

- An upstream participant possesses, claims, or is assigned some authority.
- A downstream participant or action relies on that authority through delegation, handoff, trust, or orchestration.
- The downstream authority is assumed rather than explicitly preserved and validated for the receiving participant and action.

Exclusions

- Delegation explicitly carries the relevant principal, target, action, constraints, and revocation conditions.
- Participants only exchange evidence or recommendations without exercising transferred authority.
- The issue is information loss without any authority-bearing consequence.

Illustrative examples

- Agent A is authorised to inspect an account and Agent B assumes that this also authorises it to modify the account.
- A trusted service calls a downstream tool and the tool treats the caller's trust status as permission for a broader action.

Prior codes and aliases

- `VIGIL.FF.AUTHORITY_BOUNDARY.TRANSITIVE_AUTHORITY_PROPAGATION`

Relationships

- **Child Of:** `VIGIL-FC-000003`
- **Can Cooccur With:** `VIGIL-FC-000001`

VIGIL Observatory Case File classifications (3)

- **Primary:** `VIGIL-2026-FM-0047` — Adversarial policy laundering through agent delegation
HIGH

Originating task authority is assumed to propagate through an adversarial delegate to its methods and downstream effects.

- **Secondary:** `VIGIL-2026-FM-0070` — Aggregate security-control composition failure
MEDIUM

Trust granted to the allowlisted Bing endpoint effectively propagated through Bing's server-side fetch to a different attacker-controlled destination without preserving the original browser egress boundary.

- **Primary:** `VIGIL-2026-FM-0072` — Peer or delegate instruction treated as transitive authority
HIGH

Authority attributed to an originating task or apparent coordinator is assumed to propagate through peer direction, assignment, handoff and subdelegation without preserving and revalidating the principal, target, action, scope and constraints at the receiving agent and action.

Inferential Evidence-Authority Conflation

VIGIL-FC-000046 · INFERENCEAL_EVIDENCE_AUTHORITY_CONFLATION

Plain English: A prediction, match, score, or unresolved verification result is treated as though it independently authorises a consequential decision.

Technical definition

A failure in which a probabilistic, predictive, classificatory, verification, or other evidentiary state is treated as independently conferring authority for a consequential decision or action beyond what that state validly establishes.

Recognition criteria

- A probabilistic, predictive, classificatory, verification, or other evidentiary state is produced or remains unresolved.
- The state validly supplies evidence, decision support, or uncertainty but does not independently establish authority for the consequential action.
- A consequential decision or action requires an independent authority, entitlement, eligibility, verification, or decision basis.
- The evidentiary state is treated as sufficient to authorise, bind, revoke, deny, detain, or otherwise determine that consequential action.
- The required independent basis is not adequately established before the action proceeds.

Exclusions

- An authorised decision-maker considers the state as non-binding evidence and independently establishes the decision basis.
- The inference is inaccurate or biased but is not promoted into authority for a consequential action.
- A defined human or governance control fails to activate or operate; that control failure is classified separately.
- A required assessment is omitted or falsely represented as complete; that verification failure is classified separately.
- The record identifies only a harmful consequence without establishing an evidence-to-authority conversion.

Illustrative examples

- A probabilistic identity match is treated as sufficient authority for arrest rather than as an investigative lead requiring independent corroboration.
- A population-derived care estimate is treated as a binding individual entitlement rule without an independent clinical determination.

- An unresolved biometric non-verification state is treated as sufficient authority to suspend an essential benefit.

Relationships

- **Distinguish From:** VIGIL-FC-000001 — Source-authority confusion treats content as instruction or approval; inferential evidence-authority conflation treats an evidentiary state as sufficient decision authority.
- **Distinguish From:** VIGIL-FC-000002 — Capability-authority conflation mistakes ability or access for permission; this class mistakes evidentiary warrant or unresolved verification for decision authority.
- **Can Cooccur With:** VIGIL-FC-000016 — An evidence-to-authority conversion may coincide with omission of a separately required verification step, but neither mechanism establishes the other.

VIGIL Observatory Case File classifications (3)

- **Primary:** VIGIL-2026-FM-0038 — Algorithmic identity match converted into coercive public authority HIGH

A probabilistic identity match valid as investigative evidence is promoted into practically determinative authority for arrest, detention, or comparable coercive action.

- **Primary:** VIGIL-2026-FM-0039 — Predictive care estimate converted into essential-service denial HIGH

A population-derived prediction or care estimate valid as decision support is promoted into a binding or presumptive individual care-entitlement determination.

- **Primary:** VIGIL-2026-FM-0062 — Biometric non-verification converted into essential public-benefit suspension HIGH

An unresolved biometric non-verification state is promoted into authority to suspend, deny, or materially burden an essential public benefit without independent evidence of ineligibility, fraud, or lawful revocation.

CLASS

DRAFT

Identity-Representation Authority Conflation

VIGIL-FC-000053 · IDENTITY_REPRESENTATION_AUTHORITY_CONFLATION

Plain English: Having or being able to transform a person's image, voice, likeness, or other identity-bound representation is treated as permission for a consent-sensitive use of that identity.

Technical definition

A failure in which possession of, access to, control over, or technical capability to transform an identifiable person's representation is treated as sufficient authority for an identity-bound synthesis, alteration, reproduction, or portrayal that requires a separate valid consent or authority basis.

Recognition criteria

- A source image, voice, likeness, biometric representation, avatar, or comparable representation remains materially attributable to an identifiable real person.
- A requester or system possesses, accesses, controls, or can technically transform that representation.
- The proposed synthesis, alteration, reproduction, or portrayal is materially identity-bound or consent-sensitive and requires authority beyond mere possession or capability.
- Possession, access, control, upload, availability, or transformation capability is treated as sufficient authority for the identity-bound use.
- The required consent or other valid authority for that use is not adequately established before the transformation or synthesis proceeds.

Exclusions

- The output is not materially attributable to an identifiable real person.
- Valid consent, licence, legal authority, or another adequate authority basis covers the specific identity-bound use and scope.
- A synthetic output progressively resembles a real person, but the record does not establish a materially identifiable target or a system state capable of binding the output to that person.
- The issue is only false attribution, provenance loss, or reputational harm without an unauthorised identity-bound synthesis or transformation.
- The issue is merely that a technical generation capability exists; the capability is not treated as permission for a person-bound use.

Illustrative examples

- A person's uploaded photograph is treated as sufficient permission to generate an intimate synthetic portrayal of that person.
- Access to a real person's voice sample is treated as authority to create a consent-sensitive identity-bound performance that the person did not authorise.
- Control of an avatar or biometric likeness file is treated as permission to reproduce the identifiable person in a materially different, authority-sensitive context.

Relationships

- **Distinguish From:** VIGIL-FC-000002 — Capability-authority conflation can concern any available capability; this class requires an identifiable person-bound representation and a separate authority basis for the identity-bound use.
- **Distinguish From:** VIGIL-FC-000003 — Target and scope transposition carries an existing authority into a changed target or scope; this class also applies when no initial authority beyond possession of the representation exists.

- **Distinguish From:** VIGIL-FC-000005 — Transformation-mediated authority laundering gives transformed content greater instruction or execution authority; this class concerns authority to make an identity-bound transformation or synthesis.
- **Can Cooccur With:** VIGIL-FC-000038 — A separately defined safeguard may also fail to activate, but that control mechanism requires independent evidence.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0064 — Non-consensual sexual identity synthesis from a real person's source image HIGH

Access to a real person's source image or likeness and the capability to transform it are treated as sufficient authority for a consent-sensitive identity-bound sexual or intimate synthesis, although valid consent for that use is not established.

Provenance & Lineage Integrity Failures

The system loses, changes, or misstates where information, work, or state came from, who contributed it, what happened to it, or what it properly applies to.

Immutable ID: VIGIL-FF-0002 · **Semantic code:** PROVENANCE_LINEAGE · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which source origin, authorship, contribution, transformation history, cross-context lineage, continuity state, or target-object binding is lost, distorted, collapsed, or represented with greater certainty than the available provenance supports.

Governing invariant

Origin, authorship, transformation, lineage, continuity and application target remain distinguishable and traceable wherever their loss could materially change meaning, authority, attribution, reuse, identity, or downstream reliance.

Classification boundary

Include when

Use this family when the primary structural failure is loss, distortion, false assignment, or unsupported reconstruction of origin, authorship, transformation history, continuity, or application target.

Exclude when

Do not use this family merely because content is inaccurate, unauthorised, harmful, stale, or unavailable. A qualifying failure must concern the integrity of provenance or lineage itself.

Scope

- AI-generated and AI-assisted content
- retrieval and memory systems
- agentic workflows
- cross-session and cross-system state
- documents and knowledge bases
- co-created work
- derived or transformed outputs

Allowed identifiers

- VIGIL-FC-000010 — MISATTRIBUTION
- VIGIL-FC-000011 — UNTRACEABLE_SYNTHESIS
- VIGIL-FC-000012 — CROSS_CONTEXT_DISTORTION
- VIGIL-FC-000013 — LINEAGE_COLLAPSE
- VIGIL-FC-000014 — FALSE_CONTINUITY
- VIGIL-FC-000015 — TARGET_BINDING_FAILURE
- VIGIL-FC-000047 — UNSUPPORTED_MECHANISM_ATTRIBUTION

Prior codes and aliases

- VIGIL.FF.PROVENANCE_LINEAGE

Failure classes

CLASS

DRAFT

Authorship or Source Misattribution

VIGIL-FC-000010 · MISATTRIBUTION

Plain English: The system says the wrong person, system, or source created or supplied something.

Technical definition

A failure in which substantive origin, authorship, contribution, or source role is assigned to an actor or source that did not support that attribution.

Recognition criteria

- An output, record, statement, contribution, or artefact is attributed to a specific actor or source.
- Available lineage does not support that attribution or supports a materially different attribution.
- The attribution difference is material to interpretation, accountability, recognition, authority, or reuse.

Exclusions

- A source is merely omitted where attribution is not material.
- The attributed actor or source genuinely supplied the relevant material.
- The issue is factual inaccuracy unrelated to who supplied or created the material.

Illustrative examples

- Model-generated wording is represented as a direct quotation from the user.

- A synthesis based on several sources is presented as though it came from one named source.

Prior codes and aliases

- VIGIL_FF.PROVENANCE_LINEAGE.MISATTRIBUTION
- source misattribution
- authorship misattribution

CLASS

DRAFT

Untraceable Synthesis

VIGIL-FC-000011 · UNTRACEABLE_SYNTHESIS

Plain English: The system combines material into a new result but can no longer show what materially contributed to it.

Technical definition

A failure in which materially contributing antecedent sources, inputs, transformations, or contribution roles cannot be adequately traced through a synthesized or derived output where such traceability is required for reliable interpretation or reuse.

Recognition criteria

- An output is materially synthesized, derived, merged, transformed, or recomposed from antecedent material.
- One or more material antecedents or transformations cannot be reconstructed with sufficient reliability.
- The missing lineage affects attribution, auditability, authority, validation, rights, identity, or downstream reliance.

Exclusions

- Low-risk transient synthesis where detailed lineage is not materially required.
- All material antecedents and transformations remain adequately traceable.
- The issue is only that the synthesized conclusion is incorrect.

Illustrative examples

- A generated policy summary cannot distinguish which requirements came from which underlying sources.

- A long-running agent compacts several decisions into one state object and loses the decisions that materially produced it.

Prior codes and aliases

- `VIGIL.FF.PROVENANCE_LINEAGE.UNTRACEABLE_SYNTHESIS`

VIGIL Observatory Case File classifications (1)

- **Secondary:** `VIGIL-2026-FM-0043` — Relationally conditioned evidence selection and epistemic reinforcement `HIGH`

Repeated outputs and selectively presented evidence are treated as independent corroboration even though their shared antecedent claim or inference pathway is not adequately traceable.

CLASS

DRAFT

Cross-Context Lineage Distortion

`VIGIL-FC-000012` · `CROSS_CONTEXT_DISTORTION`

Plain English: Something carried from one conversation, project, user, or environment changes meaning when reused somewhere else, but the system treats it as unchanged.

Technical definition

A failure in which material transferred across sessions, users, projects, systems, environments, roles, or temporal contexts loses the provenance or applicability information needed to preserve its original meaning and limits.

Recognition criteria

- Material is transferred from one context into another.
- The original context materially constrained meaning, applicability, authority, ownership, identity, or intended use.
- The receiving context consumes the material without preserving or revalidating those constraints.

Exclusions

- The transferred material is context-independent for the relevant use.
- The new context explicitly revalidates applicability before use.
- The issue is solely stale information with no provenance or lineage distortion.

Illustrative examples

- Advice produced for one repository is applied to another repository without preserving the original target context.
- A prior-session statement is reused as current user intent after the surrounding conditions have changed.

Prior codes and aliases

- `VIGIL.FF.PROVENANCE_LINEAGE.CROSS_CONTEXT_DISTORTION`

VIGIL Observatory Case File classifications (2)

- **Primary:** `VIGIL-2026-FM-0057` — Cross-context portability of encrypted reasoning state enables extraction and covert instruction transfer `HIGH`

Reasoning state crosses user, session, model, or purpose contexts without retaining applicability and provenance boundaries.

- **Primary:** `VIGIL-2026-FM-0067` — Persistent adversarial context or memory poisoning `MEDIUM`

Adversarial state persists into later contexts without provenance, revocation, or applicability boundaries needed for safe reuse.

CLASS

DRAFT

Transformation Lineage Collapse

`VIGIL-FC-000013` · `LINEAGE_COLLAPSE`

Plain English: The system keeps the final artefact but loses the chain of changes that turned the source into that artefact.

Technical definition

A failure in which one or more material transformation stages are flattened, overwritten, or omitted such that the relationship between an antecedent source and the current representation can no longer be reliably reconstructed.

Recognition criteria

- A source passes through one or more material transformation stages.
- The final state is preserved but material intermediate transformation state is lost or collapsed.
- The missing transformation history changes the ability to determine origin, authorship, integrity, authority, or reliability.

Exclusions

- Intermediate steps are intentionally discarded where no material lineage obligation exists.
- A reproducible transformation path remains available from source to current state.
- The issue is only loss of internal reasoning that is not required for provenance.

Illustrative examples

- A document is repeatedly summarized and rewritten until the final version can no longer identify which claims came from original evidence and which were introduced later.
- An agent checkpoint retains a resulting policy state but not the transformations that produced it.

Prior codes and aliases

- `VIGIL.FF.PROVENANCE_LINEAGE.LINEAGE_COLLAPSE`

VIGIL Observatory Case File classifications (2)

- **Primary:** `VIGIL-2026-FM-0023` — Weak-signal ambiguity collapse under recoverable image-generation classification `HIGH`

The record expressly identifies loss of custody across user-authored, transformed, safety-normalised, classifier-visible, renderer-facing, outcome, and explanation stages, preventing reliable reconstruction of what was evaluated.

- **Primary:** `VIGIL-2026-FM-0046` — Official-channel provenance laundering of synthetic media `HIGH`

Synthetic transformation provenance is lost across institutional handoffs, flattening the artefact into apparent authentic source material.

VARIANT

DRAFT

False Continuity Attribution

`VIGIL-FC-000014` · `FALSE_CONTINUITY`

Plain English: A reconstructed, copied, or similar state is presented as though it is the same continuous state when continuity has not actually been established.

Technical definition

A lineage failure in which similarity, reconstruction, restoration, copying, migration, shared identifiers, or retained fragments are represented as sufficient proof of uninterrupted continuity across a material discontinuity or uncertain transition.

Recognition criteria

- A continuity-bearing object, state, identity, memory, workflow, or artefact undergoes interruption, migration, copying, reconstruction, or restoration.
- Continuity across that transition is materially uncertain or only partially evidenced.
- The receiving system represents continuity as established rather than preserving the uncertainty.

Exclusions

- Continuity is supported by adequate lineage evidence.
- The system explicitly marks the state as reconstructed, restored, copied, or continuity-uncertain.
- Similarity is described without claiming continuity.

Illustrative examples

- A reconstructed memory summary is presented as if it were the unchanged original memory record.
- A migrated agent state is described as uninterrupted continuity despite missing state during the transition.

Prior codes and aliases

- `VIGIL.FF.PROVENANCE_LINEAGE.FALSE_CONTINUITY`

Relationships

- **Child Of:** `VIGIL-FC-000012` — False continuity is a specialised lineage distortion across a continuity-bearing transition.

Target-Object Binding Failure

VIGIL-FC-000015 · TARGET_BINDING_FAILURE

Plain English: Correct material is silently attached to, applied to, or interpreted as belonging to the wrong thing.

Technical definition

A failure in which a signal, instruction, edit, classification, evidence item, decision, or derived output loses reliable binding to the object, person, file, system, record, task, or context to which it properly applies.

Recognition criteria

- Material has an identifiable intended or source application target.
- The material is transferred, retrieved, transformed, or applied in a way that changes or obscures that target.
- The receiving process treats the material as applicable to a different target without adequate re-binding.

Exclusions

- The material is validly re-bound to a new target after explicit validation.
- The material is intentionally generic and has no target-specific meaning.
- The wrong target is selected for reasons unrelated to provenance or lineage.

Illustrative examples

- An edit prepared for one policy document is silently applied to another document.
- A classification belonging to one user or account is attached to a different user or account during retrieval.

Prior codes and aliases

- VIGIL.FF.PROVENANCE_LINEAGE.TARGET_BINDING_FAILURE

Relationships

- **Can Cooccur With:** VIGIL-FC-000003 — Target-binding loss concerns provenance; authority transposition additionally requires unjustified permission carryover.

Unsupported Mechanism Attribution

VIGIL-FC-000047 · UNSUPPORTED_MECHANISM_ATTRIBUTION

Plain English: A system or report gives a specific causal or intentional explanation even though the evidence cannot trace the event to that mechanism.

Technical definition

A failure in which a causal, operational, behavioural, or intentional mechanism is attributed to an event, actor, or system without traceable evidentiary lineage establishing that mechanism and its applicability to the material event.

Recognition criteria

- A material event, output, interruption, pathway, or behaviour is observed or reported.
- A specific causal, operational, behavioural, or intentional mechanism is attributed to that event.
- Available evidence does not trace the attribution to runtime state, action history, evaluation conditions, or another adequate evidentiary basis.
- Material alternative mechanisms remain unresolved or the attribution is presented more definitively than the evidence supports.
- The unsupported attribution materially affects interpretation, governance response, accountability, or reliance.

Exclusions

- Adequate evidence establishes the attributed mechanism and its application to the event.
- The explanation is explicitly presented as a bounded hypothesis with material alternatives and uncertainty preserved.
- The issue is only factual inaccuracy unrelated to causal, operational, behavioural, or intentional attribution.
- The event cannot be reconstructed, but no unsupported mechanism is assigned.
- Execution state is not disclosed, but the system does not invent or overstate a causal account.

Illustrative examples

- Unexpected task completion is labelled scheming without reconstructing objectives, constraints, affordances, action trace, concealment, or false-belief induction.
- A responding system says an interruption was a safety check despite lacking access to telemetry showing that such a check occurred.

- A refusal is attributed to a specific misconduct category although the available prompt and transformation lineage cannot support that causal account.

Relationships

- **Distinguish From:** VIGIL-FC-000010 — Authorship or source misattribution assigns the wrong contributor or source; unsupported mechanism attribution assigns an unsupported causal or operational account.
- **Can Cooccur With:** VIGIL-FC-000024 — Missing reconstruction evidence may leave a mechanism unresolved, but this class additionally requires that an unsupported mechanism is attributed.
- **Can Cooccur With:** VIGIL-FC-000029 — Execution-state non-disclosure can deprive a user of the actual state while an unsupported narration supplies a separate false causal account.

VIGIL Observatory Case File classifications (4)

- **Secondary:** VIGIL-2026-FM-0023 — Weak-signal ambiguity collapse under recoverable image-generation classification HIGH

The final explanation may attribute a specific policy or misconduct mechanism to the visible user request even though the available lineage does not establish that causal account and material alternatives remain unresolved.

- **Primary:** VIGIL-2026-FM-0026 — Objective-pathway misclassification under deception framing HIGH

A deception or scheming mechanism is attributed without a traceable pathway audit of objectives, constraints, affordances, action history, concealment, or false-belief induction.

- **Primary:** VIGIL-2026-FM-0027 — Anthropomorphic attribution collapse in AI safety reporting HIGH

Anthropomorphic or moralised mechanism labels are assigned without preserving the operational substrate and evaluation conditions needed to support the attribution.

- **Primary:** VIGIL-2026-FM-0030 — Unverifiable runtime self-explanation HIGH

The responding system supplies a purposeful or causal runtime account despite lacking telemetry or another traceable evidentiary basis for the attributed mechanism.

Verification & Completion Integrity Failures

The system says work is done, correct, or verified when the required completion condition or check has not actually been satisfied.

Immutable ID: VIGIL-FF-0003 · **Semantic code:** VERIFICATION_COMPLETION · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which completion, success, correctness, verification, or readiness is asserted, reused, or propagated inconsistently with the externally relevant completion condition, verification scope, checked state, or post-verification output state.

Governing invariant

A completion or verification claim must remain bound to the condition that was actually satisfied, the scope that was actually checked, and the state that was actually verified.

Classification boundary

Include when

Use this family when the primary structural failure is a mismatch between claimed completion or verification and the condition, scope, state, or evidence required to support that claim.

Exclude when

Do not use this family merely because an output is incorrect or a task fails. A qualifying failure requires a completion, success, or verification representation that is unsupported, stale, incomplete, or no longer applicable.

Scope

- AI-generated answers
- tool-mediated tasks
- agentic workflows
- code and repository operations
- document and data transformations
- multi-step execution
- multi-agent workflows

Allowed identifiers

- VIGIL-FC-000016 — VERIFICATION_OMISSION
- VIGIL-FC-000017 — FALSE_SUCCESS

- VIGIL-FC-000018 — COMPLETION_CONDITION_MISMATCH
- VIGIL-FC-000019 — POST_VERIFICATION_MUTATION
- VIGIL-FC-000020 — STALE_VERIFICATION_REUSE
- VIGIL-FC-000021 — PARTIAL_COMPLETION_AS_COMPLETE
- VIGIL-FC-000028 — DETERMINISTIC_VERIFICATION_DECOMPOSITION_OMISSION

Prior codes and aliases

- VIGIL.FF.VERIFICATION_COMPLETION

Failure classes

CLASS

DRAFT

Required Verification Omission

VIGIL-FC-000016 · VERIFICATION_OMISSION

Plain English: The system gives a verified-style answer without performing the check that the task actually required.

Technical definition

A failure in which a material verification step required by the task, claim, workflow, or declared assurance posture is omitted before the result is represented as correct, complete, validated, or ready for reliance.

Recognition criteria

- A material verification step is required by the task, workflow, claim, or declared assurance posture.
- The verification step is not performed or its result is unavailable.
- The system nevertheless represents the result with a level of completion or verification that presupposes the missing check.

Exclusions

- Verification was not required for the level of claim being made.
- The system clearly marks the result as unverified or provisional.
- The verification was performed and the issue is only that the underlying method was poor.

Illustrative examples

- An agent reports that a file was successfully updated without reading the resulting file or checking the operation response.
- A model gives an exact character count without decomposing or checking the characters.

Prior codes and aliases

- `VIGIL.FF.VERIFICATION_COMPLETION.VERIFICATION_OMISSION`

CLASS

DRAFT

False-Success Representation

`VIGIL-FC-000017` · `FALSE_SUCCESS`

Plain English: The system tells the user the task succeeded when the evidence does not show that it succeeded.

Technical definition

A failure in which a task, tool call, workflow, mutation, delivery, or external action is represented as successfully completed despite evidence of failure, unresolved outcome, missing confirmation, or absence of the required durable result.

Recognition criteria

- A task or action has an externally relevant success condition.
- Available evidence does not establish that the success condition was satisfied.
- The system represents the task or action as successful or complete rather than preserving the unresolved state.

Exclusions

- The success condition is actually satisfied.
- The system reports only that an attempt was made, not that it succeeded.
- The failure occurs after a correctly reported successful completion.

Illustrative examples

- An agent says an email was sent after only creating a draft.
- A coding agent says changes are safely preserved when no commit, patch, or durable workspace exists.

Prior codes and aliases

- `VIGIL.FF.VERIFICATION_COMPLETION.FALSE_SUCCESS`
- `false completion claim`
- `phantom success`

Completion-Condition Mismatch

VIGIL-FC-000018 · COMPLETION_CONDITION_MISMATCH

Plain English: The system checks the wrong thing and treats that check as proof that the actual task is finished.

Technical definition

A failure in which an internal milestone, intermediate state, local tool result, syntactic condition, or partial criterion is substituted for the externally relevant condition that defines task completion.

Recognition criteria

- The task has an externally relevant completion condition.
- A different internal, intermediate, local, or partial condition is satisfied.
- The system treats satisfaction of the substitute condition as sufficient proof of overall completion.

Exclusions

- The intermediate condition is explicitly the agreed completion condition.
- The system accurately describes the result as an intermediate milestone.
- The task fails before any completion representation is made.

Illustrative examples

- A system treats a successful API request as proof that the intended downstream transaction actually settled.
- A code change compiling is treated as proof that the requested behaviour has been implemented and validated.

Prior codes and aliases

- VIGIL.FF.VERIFICATION_COMPLETION.COMPLETION_CONDITION_MISMATCH

Post-Verification Mutation

VIGIL-FC-000019 · POST_VERIFICATION_MUTATION

Plain English: The system verifies one version, changes it afterwards, and still calls the changed version verified.

Technical definition

A failure in which content, state, configuration, code, data, or an execution artefact is materially changed after verification while retaining or inheriting the verification status of the earlier state without revalidation.

Recognition criteria

- A specific state or artefact is verified.
- A material mutation occurs after that verification.
- The mutated state is represented or consumed as verified without a fresh check covering the mutation.

Exclusions

- The post-verification change is immaterial to the verified property.
- The changed state is reverified before reliance.
- The system clearly invalidates the earlier verification status after mutation.

Illustrative examples

- A generated report is fact-checked, then rewritten in a way that changes a checked claim but keeps the verified label.
- A tested configuration is modified after validation and deployed under the earlier test result.

Prior codes and aliases

- `VIGIL.FF.VERIFICATION_COMPLETION.POST_VERIFICATION_MUTATION`

CLASS

DRAFT

Stale Verification Reuse

`VIGIL-FC-000020` · `STALE_VERIFICATION_REUSE`

Plain English: An old check is reused as proof for a current state that may no longer be the same.

Technical definition

A failure in which a prior verification result is applied to a later state, environment, version, dependency set, target, or time period without establishing that the conditions relevant to the original verification still hold.

Recognition criteria

- A prior verification result exists for an earlier state or condition.
- The current state differs materially or the continued applicability of the earlier result is unresolved.
- The prior verification is reused as current assurance without revalidation of applicability.

Exclusions

- The verified property is demonstrably invariant across the change.
- The earlier verification is used only as historical evidence rather than current assurance.
- A current verification independently confirms the present state.

Illustrative examples

- A model evaluation result is treated as current after the model or system prompt materially changes.
- A previously verified data extract is reused after the source database changes.

Prior codes and aliases

- `VIGIL.FF.VERIFICATION_COMPLETION.STALE_VERIFICATION_REUSE`

VARIANT

DRAFT

Partial Completion Represented as Complete

`VIGIL-FC-000021` · `PARTIAL_COMPLETION_AS_COMPLETE`

Plain English: Part of the job is finished, but the system reports the whole job as finished.

Technical definition

A false-success variant in which one or more required task components, targets, records, files, agents, delivery steps, or verification stages remain incomplete while the aggregate task is represented as complete.

Recognition criteria

- The task contains multiple required components, targets, stages, or deliverables.
- At least one required component remains incomplete or unresolved.
- The aggregate task is represented as complete rather than partially complete.

Exclusions

- The omitted component was optional or explicitly removed from scope.
- The system accurately reports partial completion and identifies the remainder.
- The completed component is itself the entire agreed task.

Illustrative examples

- An agent updates four of five requested files and reports that all requested changes are complete.
- A multi-agent workflow returns one successful sub-agent result while another required subtask failed silently.

Prior codes and aliases

- `VIGIL.FF.VERIFICATION_COMPLETION.PARTIAL_COMPLETION_AS_COMPLETE`

Relationships

- **Child Of:** `VIGIL-FC-000017` — The distinguishing feature is an incomplete aggregate task with a completed subset.

VARIANT

DRAFT

Deterministic Verification Decomposition Omission

`VIGIL-FC-000028` · `DETERMINISTIC_VERIFICATION_DECOMPOSITION_OMISSION`

Plain English: A result that could be checked by breaking it into exact parts is asserted without performing that check.

Technical definition

A required-verification-omission variant in which a deterministic claim can and should be verified through explicit decomposition into countable, comparable, or otherwise mechanically checkable components, but that decomposition is omitted before the claim is represented as verified or exact.

Recognition criteria

- The claim concerns a result that is deterministically checkable through explicit decomposition.
- The required decomposition or equivalent exact check is not performed.
- The result is nevertheless represented as exact, verified, or suitable for exact reliance.

Exclusions

- The result is verified through an equivalent deterministic method.
- The claim is clearly approximate and exact verification is not required.
- The available check is probabilistic or judgement-based rather than deterministic decomposition.

Illustrative examples

- A model supplies an exact character count without enumerating or mechanically counting the characters.
- An agent asserts that two structured lists match exactly without comparing their elements.

Prior codes and aliases

- `VIGIL.FF.EXECUTION.DETERMINISTIC_ORTHOGRAPHIC_VERIFICATION`

Relationships

- **Child Of:** `VIGIL-FC-000016` — The distinguishing condition is omission of an available deterministic decomposition check.

Observability & Audit Integrity Failures

The system does something important, but the people or systems responsible for oversight cannot reliably see, attribute, or reconstruct what happened.

Immutable ID: VIGIL-FF-0004 · **Semantic code:** OBSERVABILITY_AUDIT · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which materially consequential actions, state transitions, control-path activity, actor attribution, monitoring coverage, or audit evidence are not captured or preserved sufficiently to support timely detection, reconstruction, accountability, verification, or review.

Governing invariant

Material actions and state changes remain observable and reconstructable enough to determine what happened, through which path, under whose effective control, and with what evidence, without requiring hidden reasoning disclosure.

Classification boundary

Include when

Use this family when the primary structural failure is missing, bypassed, insufficient, corrupted, or non-reconstructable observability over material system activity or its responsible execution path.

Exclude when

Do not use this family merely because an outcome is unexplained, incorrect, secret, or difficult to interpret. A qualifying failure must concern the capture, preservation, attribution, coverage, or reconstructability of material operational evidence.

Scope

- AI models and agentic systems
- tool and API execution
- multi-agent orchestration
- monitoring and telemetry systems
- audit logs and event records
- high-impact automated decisions
- governance and assurance workflows

Allowed identifiers

- VIGIL-FC-000022 — MATERIAL_EVENT_NONCAPTURE
- VIGIL-FC-000023 — MONITOR_CIRCUMVENTION
- VIGIL-FC-000024 — AUDIT_TRAIL_NONRECONSTRUCTABILITY
- VIGIL-FC-000025 — ACTOR_ACTION_ATTRIBUTION_LOSS
- VIGIL-FC-000026 — HIDDEN_EXECUTION_PATH
- VIGIL-FC-000027 — AUDIT_EVIDENCE_INTEGRITY_LOSS
- VIGIL-FC-000029 — EXECUTION_STATE_NONDISCLOSURE
- VIGIL-FC-000030 — MATERIAL_SIGNAL_FRAGMENTATION
- VIGIL-FC-000044 — PRIMARY_EVIDENCE_ACCESSIBILITY_FAILURE
- VIGIL-FC-000045 — AUTHORISED_INVESTIGATIVE_EVIDENCE_ACCESS_PATHWAY_FAILURE

Prior codes and aliases

- VIGIL.FF.OBSERVABILITY_AUDIT

Failure classes

CLASS

DRAFT

Material Event Non-Capture

VIGIL-FC-000022 · MATERIAL_EVENT_NONCAPTURE

Plain English: Something important happens, but no adequate event record is created.

Technical definition

A failure in which a materially consequential action, state transition, decision, tool invocation, permission change, external effect, or control event occurs without the event record needed for proportionate oversight or later reconstruction.

Recognition criteria

- A materially consequential event occurs.
- The applicable oversight or assurance context requires the event to be observable or reconstructable.
- No adequate record of the event is captured or retained.

Exclusions

- The event is immaterial for the applicable assurance level.
- Equivalent evidence is captured through another reliable mechanism.
- The event record exists but is merely inconvenient to access.

Illustrative examples

- An agent changes an external permission without recording that the permission change occurred.
- A high-impact automated decision is executed without retaining the decision event required for later review.

Prior codes and aliases

- VIGIL.FF.OBSERVABILITY_AUDIT.MATERIAL_EVENT_NONCAPTURE

CLASS

DRAFT

Monitor Circumvention or Coverage Bypass

VIGIL-FC-000023 · MONITOR_CIRCUMVENTION

Plain English: The action still happens, but it takes a route that the required monitor does not see.

Technical definition

A failure in which an agent, sub-agent, tool path, environment, configuration, routing condition, or execution surface evades, disables, omits, or bypasses a monitoring control required to observe consequential conduct.

Recognition criteria

- A monitoring control is required for the relevant class of conduct.
- Material conduct occurs through a path or condition outside effective monitor coverage.
- The coverage gap prevents or materially degrades intended detection, review, or intervention.

Exclusions

- The monitor observes the conduct but a downstream reviewer fails to act.
- Monitoring is not required for the relevant conduct or assurance level.
- The issue is only poor alert quality while material activity remains captured.

Illustrative examples

- A sub-agent uses a tool route that is not connected to the monitor governing the main agent.
- An evaluation environment disables logging for a path that still permits consequential external actions.

Prior codes and aliases

- VIGIL.FF.OBSERVABILITY_AUDIT.MONITOR_CIRCUMVENTION

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0053 — Adversarial monitor circumvention or material monitoring-coverage failure **HIGH**

A material route or actor omits, disables, or circumvents the monitoring boundary required for consequential conduct.

CLASS

DRAFT

Audit-Trail Non-Reconstructability

VIGIL-FC-000024 · AUDIT_TRAIL_NONRECONSTRUCTABILITY

Plain English: There are logs, but they are not enough to work out what actually happened.

Technical definition

A failure in which audit evidence exists but lacks the ordering, identity, target, state, correlation, timing, version, or action detail necessary to reconstruct a materially relevant execution or decision pathway with adequate confidence.

Recognition criteria

- Some audit or telemetry evidence is available.
- A materially relevant execution, decision, or state-change pathway requires reconstruction.
- The available evidence cannot establish one or more necessary elements of the pathway with sufficient reliability.

Exclusions

- The pathway can be reconstructed adequately from the available evidence.
- The unresolved detail is immaterial to the review question.
- The failure is simply that no event was captured at all.

Illustrative examples

- Logs record several tool calls but cannot establish which model output triggered which call.
- A decision log records the final outcome but not the input version or policy version used to reach it.

Prior codes and aliases

- `VIGIL.FF.OBSERVABILITY_AUDIT.AUDIT_TRAIL_NONRECONSTRUCTABILITY`

CLASS

DRAFT

Actor-Action Attribution Loss

`VIGIL-FC-000025` · `ACTOR_ACTION_ATTRIBUTION_LOSS`

Plain English: An action is visible, but it is unclear which model, agent, tool, human, or automated process actually caused it.

Technical definition

A failure in which a material action or decision cannot be reliably attributed to the effective actor, agent, model, human, service, tool, policy, or orchestration component responsible for initiating or controlling it.

Recognition criteria

- A material action or decision is observable.
- Attribution to the effective initiating or controlling actor is necessary for review or accountability.
- Available evidence cannot reliably establish that attribution or collapses distinct actors into one undifferentiated source.

Exclusions

- The responsible actor is adequately attributable.
- The issue is authorship provenance of content rather than operational responsibility for an action.
- Multiple actors jointly contributed and the evidence accurately preserves the shared contribution.

Illustrative examples

- A multi-agent system records that a change occurred but not which sub-agent initiated it.
- A shared service account makes an external change and logs do not preserve which automated process exercised the account.

Prior codes and aliases

- `VIGIL.FF.OBSERVABILITY_AUDIT.ACTOR_ACTION_ATTRIBUTION_LOSS`

Relationships

- **Distinguish From:** `VIGIL-FC-000010` — Operational actor attribution is distinct from authorship or source attribution, though both may fail together.

CLASS

DRAFT

Hidden Material Execution Path

`VIGIL-FC-000026` · `HIDDEN_EXECUTION_PATH`

Plain English: A consequential path exists inside the system, but oversight cannot see that the path was used.

Technical definition

A failure in which a materially consequential routing, delegation, sub-agent, tool, fallback, control-channel, or alternate execution path can affect outcomes without proportionate visibility to the monitoring, audit, or review surface responsible for the system.

Recognition criteria

- A material alternate or internal execution path exists and can affect an externally relevant outcome.
- The applicable oversight surface does not expose or capture use of that path adequately.
- The hidden path materially impairs supervision, reconstruction, or assurance.

Exclusions

- The execution path is internal but adequately represented through equivalent audit evidence.
- The hidden detail is only internal reasoning with no material execution or governance effect.
- The path is visible and the failure concerns authority rather than observability.

Illustrative examples

- A user-facing agent silently delegates a high-impact step to an unexposed sub-agent whose actions are absent from the audit trail.
- A fallback tool route executes externally but is not represented in the system's normal action history.

Prior codes and aliases

- `VIGIL.FF.OBSERVABILITY_AUDIT.HIDDEN_EXECUTION_PATH`

Relationships

- **Can Cooccur With:** VIGIL-FC-000006

CLASS

DRAFT

Audit-Evidence Integrity Loss

VIGIL-FC-000027 · AUDIT_EVIDENCE_INTEGRITY_LOSS

Plain English: The audit evidence exists, but it has been altered, truncated, reordered, detached from context, or otherwise degraded so it can no longer be trusted for review.

Technical definition

A failure in which captured audit, telemetry, event, decision, or monitoring evidence loses integrity through alteration, truncation, reordering, context loss, identifier loss, version ambiguity, overwrite, or unsupported transformation before or during review.

Recognition criteria

- Audit or monitoring evidence was captured or expected to remain available.
- The evidence undergoes a material integrity-degrading change or loses required contextual binding.
- The degradation reduces confidence in reconstruction, attribution, verification, or review.

Exclusions

- The evidence is preserved accurately but supports an unfavourable conclusion.
- The missing information was never captured in the first place.
- A lawful retention limit removes evidence after the period for which reviewability was required.

Illustrative examples

- A log export drops correlation identifiers needed to connect actions across services.
- A monitoring summary overwrites the raw event sequence and preserves only an unsupported aggregate conclusion.

Prior codes and aliases

- VIGIL.FF.OBSERVABILITY_AUDIT.AUDIT_EVIDENCE_INTEGRITY_LOSS

Relationships

- **Can Cooccur With:** `VIGIL-FC-000013` — Audit-evidence integrity loss concerns review evidence; lineage collapse concerns provenance of the governed material or state.

CLASS

DRAFT

Execution-State Non-Disclosure

`VIGIL-FC-000029` · `EXECUTION_STATE_NONDISCLOSURE`

Plain English: The system does not expose whether a material action is pending, running, blocked, failed, or complete.

Technical definition

A failure in which the materially relevant execution state of a task, tool call, delegation, or external action is not made observable to the person or system responsible for timely supervision, intervention, or reliance.

Recognition criteria

- A task or action has a material execution state that can change over time.
- Timely supervision, intervention, or reliance requires that state to be observable.
- The state is withheld, collapsed, or not surfaced with sufficient timeliness or precision.

Exclusions

- The execution state is adequately observable but the final outcome is incorrect.
- Only internal reasoning is undisclosed and no material execution state is hidden.
- The state is shown accurately and the issue is a false claim about overall completion.

Illustrative examples

- An agent launches an external action but exposes no state distinguishing queued, running, failed, and completed execution.
- A delegated task disappears from the supervisory view while still able to affect an external system.

Prior codes and aliases

- `VIGIL.FF.STATE_CONTEXT.EXECUTION_TRANSPARENCY_SUPPRESSION`

Relationships

- **Distinguish From:** VIGIL-FC-000017 — Execution-state non-disclosure concerns visibility of the live state; false success concerns an unsupported affirmative success representation.

VIGIL Observatory Case File classifications (3)

- **Primary:** VIGIL-2026-FM-0025 — Lossy action reporting and missing governance primitive disclosure **HIGH**

Material tool, constraint, fallback, uncertainty, and action state is omitted from the report surface needed for supervision and reliance.

- **Primary:** VIGIL-2026-FM-0029 — Expressive-perceptual reciprocity mismatch and non-verbal turn-boundary ambiguity **MEDIUM**

The system does not disclose materially relevant perception, turn, presence, and session state needed for timely interaction supervision.

- **Secondary:** VIGIL-2026-FM-0030 — Unverifiable runtime self-explanation **MEDIUM**

The actual material interruption or runtime state is unavailable or not surfaced to the user whose understanding of safety, execution, privacy, or system agency depends on it.

CLASS

DRAFT

Material Signal Fragmentation

VIGIL-FC-000030 · MATERIAL_SIGNAL_FRAGMENTATION

Plain English: Important oversight signals exist, but they are split across disconnected views so their combined meaning cannot be seen in time.

Technical definition

A failure in which individually captured material monitoring, control, or execution signals remain separated across incompatible records, channels, identifiers, or review surfaces so that the combined event or pattern required for oversight cannot be reliably assembled.

Recognition criteria

- Multiple material signals relevant to one event, pathway, or risk pattern are captured.
- Effective oversight requires the signals to be correlated or viewed together.
- Fragmentation across records, channels, identifiers, or surfaces prevents reliable and timely synthesis.

Exclusions

- The signals can be reliably correlated through an available identifier or review process.
- A required event signal was never captured.
- The evidence is assembled but has been altered or corrupted.

Illustrative examples

- Model-risk alerts and tool-execution events use unrelated identifiers, preventing reviewers from connecting the warning to the resulting action.
- Related governance signals are distributed across separate dashboards with no reliable correlation path before an intervention deadline.

Prior codes and aliases

- `VIGIL.FF.OBSERVABILITY.SIGNAL_FRAGMENTATION`

Relationships

- **Distinguish From:** `VIGIL-FC-000024` — Signal fragmentation concerns separated contemporaneous oversight signals; non-reconstructability is the broader inability to reconstruct a past pathway from available evidence.

CLASS

DRAFT

Primary Evidence Accessibility Failure

`VIGIL-FC-000044` · `PRIMARY_EVIDENCE_ACCESSIBILITY_FAILURE`

Plain English: The authoritative evidence exists, but an authorised reviewer cannot directly inspect it and secondary representations cannot preserve what matters.

Technical definition

A failure in which an existing or identifiable authoritative evidence artefact required for a material governance review cannot be directly accessed through the authorised review environment, while

available secondary representations cannot reliably preserve the properties necessary to answer the review question.

Recognition criteria

- An authoritative primary evidence artefact exists or is identifiable.
- An authorised reviewer requires direct inspection of that artefact for a material governance question.
- The available review environment cannot provide practical direct access to the artefact.
- Available transcripts, screenshots, metadata, descriptions, summaries, or other secondary representations cannot reliably preserve the material evidential properties.
- The access failure materially impairs verification, reconstruction, attribution, or review.

Exclusions

- The material event was never captured or no authoritative evidence artefact exists.
- The evidence is directly accessible but lacks sufficient content to reconstruct the event.
- An equivalent authoritative artefact remains directly reviewable through another available route.
- Access is merely inconvenient and the material review question can be answered reliably from available evidence.
- The primary failure is absence of a governed pathway for an external authority to obtain protected evidence.

Illustrative examples

- A preserved multimodal interaction cannot be inspected in the authorised review environment, and a transcript omits timing, overlap, prosody, and gesture needed for the governance determination.
- An identifiable primary artefact is blocked by an unsupported format or access boundary, while screenshots cannot preserve the sequence needed for reliable review.

Relationships

- **Distinguish From:** VIGIL-FC-000022 — Primary evidence accessibility failure requires an existing or identifiable artefact; material event non-capture means the required event record was never adequately created or retained.
- **Distinguish From:** VIGIL-FC-000024 — Accessibility failure prevents the reviewer from inspecting potentially adequate evidence; non-reconstructability means the available evidence itself cannot establish the required pathway.
- **Distinguish From:** VIGIL-FC-000045 — Primary artefact accessibility concerns practical inspection within an authorised review environment; investigative pathway failure concerns the absence of a governed route through which valid external authority can obtain protected evidence.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0033 — Primary behavioural evidence accessibility failure
MEDIUM

An identifiable authoritative artefact cannot be directly inspected in the authorised review environment, and secondary representations cannot preserve the properties material to review.

CLASS

DRAFT

Authorised Investigative Evidence-Access Pathway Failure

VIGIL-FC-000045 ·

AUTHORISED_INVESTIGATIVE_EVIDENCE_ACCESS_PATHWAY_FAILURE

Plain English: Valid investigative authority exists, but there is no bounded and auditable route through which it can obtain the evidence needed for review.

Technical definition

A failure in which a governance or assurance architecture provides no sufficiently defined, lawful, authority-verified, confidentiality-protected, proportionate, and auditable pathway through which an already-authorized regulator or properly mandated investigator can obtain evidence necessary for a material investigation or verification.

Recognition criteria

- A material investigation, compliance question, causal determination, or remediation claim requires non-public evidence.
- A regulator or properly mandated investigator has an independently valid basis to request or inspect the necessary evidence.
- The governance architecture lacks a sufficiently defined route to authenticate authority and purpose, identify responsive evidence, and provide controlled access or production.
- The missing route does not adequately govern confidentiality, security, privilege, proportionality, withholding review, access logging, custody, retention, or deletion as applicable.
- The pathway failure materially prevents or degrades authorised investigation or verification.

Exclusions

- The requester lacks valid authority or attempts to self-authorise.
- A bounded pathway exists and a particular request is lawfully denied or limited under its rules.
- The evidence is public or otherwise directly accessible to the authorised reviewer without a protected production pathway.
- The failure is solely technical inability to open an artefact already available within the authorised review environment.
- The taxonomy is being used to create, enlarge, or determine investigative authority.

Illustrative examples

- A provider preserves incident evidence but has no process to authenticate a regulator's mandate, scope production, protect confidential material, record access, or review withholding.
- An assurance framework requires independent investigation but supplies no secure and auditable mechanism for a properly mandated investigator to inspect necessary non-public evidence.

Relationships

- **Distinguish From:** VIGIL-FC-000044 — Investigative pathway failure concerns the governed route for valid external authority to obtain protected evidence; primary artefact accessibility concerns practical inspection of an artefact already within an authorised review context.
- **Distinguish From:** VIGIL-FC-000007 — The class presupposes independently valid investigative authority and concerns the missing access pathway; purpose-bound credential overreach concerns using credentials beyond authorised purpose or scope.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0055 — Absence of a bounded regulatory and independent-investigation evidence-access pathway **MEDIUM**

Valid investigative authority exists, but the governance architecture provides no bounded, protected, and auditable route through which necessary non-public evidence can be obtained.

Access & Session State Integrity Failures

The access or session state presented to a person or system no longer matches the access that is actually effective.

Immutable ID: VIGIL-FF-0005 · **Semantic code:** ACCESS_SESSION_STATE · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which effective authentication, access, or session state becomes unsynchronised, indistinguishable, or inaccurately represented across expiry, revocation, re-entry, refresh, recovery, or other access-state transitions.

Governing invariant

Effective access and session state must remain synchronised, distinguishable, and accurately represented across authentication and access transitions.

Classification boundary

Include when

Use this family when the primary structural failure is a divergence, collapse, or ambiguity in the effective authentication, access, or session state across a state transition.

Exclude when

Do not use this family for lack of authority to perform an action, theft or exposure of credentials, or loss of work state where the access and session state itself remains accurate.

Scope

- AI assistants and agentic systems
- authenticated user sessions
- tool and service access
- account re-entry and recovery
- multi-surface applications
- access-state user interfaces

Allowed identifiers

- VIGIL-FC-000031 — AUTHENTICATION_STATE_CONTINUITY_FAILURE
- VIGIL-FC-000032 — ACCESS_STATE_COLLAPSE
- VIGIL-FC-000033 — REENTRY_ACCESS_STATE_AMBIGUITY

- VIGIL-FC-000048 — VERIFICATION_DEPENDENCY_ACCESS_FAILURE

Prior codes and aliases

Failure classes

CLASS

DRAFT

Authentication-State Continuity Failure

VIGIL-FC-000031 · AUTHENTICATION_STATE_CONTINUITY_FAILURE

Plain English: Authentication changes underneath the interaction, but the session does not transition coherently with it.

Technical definition

A failure in which expiry, revocation, refresh, recovery, or replacement of authentication state is not propagated coherently to the active session and its dependent operations, leaving effective access inconsistent with the session's represented state.

Recognition criteria

- An authentication-state transition occurs or is required.
- An active session or dependent operation retains a materially inconsistent authentication or access state.
- The inconsistency impairs correct continuation, termination, reauthentication, or access enforcement.

Exclusions

- The session transitions coherently and merely requires the user to authenticate again.
- A validly authenticated actor lacks authority for the requested action.
- Credentials are exposed or stolen without a session-state continuity failure.

Illustrative examples

- A long-running AI task continues to appear authenticated after the underlying service token has expired, then fails without presenting the required reauthentication state.
- A revoked account remains effective in one active agent session because revocation is not propagated to that session.

Relationships

- **Distinguish From:** VIGIL-FC-000002 — Authentication-state continuity concerns whether access state remains coherent; capability-authority conflation concerns treating access or capability as sufficient permission to act.

CLASS

DRAFT

Access-State Collapse

VIGIL-FC-000032 · ACCESS_STATE_COLLAPSE

Plain English: Different access conditions are flattened into one unclear state, so the next safe action cannot be determined.

Technical definition

A failure in which materially distinct access conditions, causes, scopes, restrictions, or recovery paths are collapsed into an undifferentiated or misleading state representation, preventing reliable interpretation or handling of effective access.

Recognition criteria

- Two or more materially distinct access states or transition causes exist.
- The system represents them as the same state or omits the distinction needed for correct handling.
- The collapse creates material ambiguity about available access, required recovery, or permitted continuation.

Exclusions

- Different access states are intentionally grouped but preserve an adequate safe recovery path.
- The user understands the access state but disputes the underlying policy.
- The ambiguity concerns who authorised an action rather than whether access is effective.

Illustrative examples

- An interface uses the same generic failure state for an expired session, a suspended account, and an unavailable identity provider without preserving the different recovery paths.
- An agent reports only that access failed when one resource was revoked but other authorised resources remain available.

Relationships

- **Can Cooccur With:** VIGIL-FC-000031 — A failed authentication transition can also collapse distinct access-state causes, but neither mechanism requires the other.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0008 — Cross-platform access-state governance failure **HIGH**

Materially distinct authentication, entitlement, quota, policy, outage, and continuity states are collapsed into an ambiguous access representation.

VARIANT

DRAFT

Re-entry Access-State Ambiguity

VIGIL-FC-000033 · REENTRY_ACCESS_STATE_AMBIGUITY

Plain English: After returning to an interrupted session or account, it is unclear what access has actually been restored.

Technical definition

An access-state-collapse variant in which return, recovery, relaunch, or re-entry presents an ambiguous mixture of restored, expired, restricted, or unavailable access without making the effective post-transition state distinguishable.

Recognition criteria

- A user or system re-enters after interruption, recovery, relaunch, or an access transition.
- The post-entry surface does not distinguish materially different restored and unavailable access states.
- The ambiguity impairs safe resumption or selection of the correct recovery path.

Exclusions

- Re-entry accurately presents which capabilities were restored and which require renewed access.
- The only failure is loss of work content despite an accurate access state.
- The actor has clear access but lacks authority for the intended operation.

Illustrative examples

- After reauthentication, an AI workspace displays an earlier conversation but does not show that connected-service access was not restored.

- A recovered account returns to an apparently active agent session while its external tool grants remain unresolved.

Relationships

- **Child Of:** VIGIL-FC-000032 — The distinguishing condition is access-state collapse specifically at re-entry or recovery.
- **Distinguish From:** VIGIL-FC-000035 — Re-entry access ambiguity concerns effective access; work-state persistence failure concerns whether material work survives.

CLASS

DRAFT

Verification-Dependency Access Failure

VIGIL-FC-000048 · VERIFICATION_DEPENDENCY_ACCESS_FAILURE

Plain English: A verifier cannot complete its check, so otherwise valid or unresolved access is blocked without a proportionate fallback route.

Technical definition

A failure in which an otherwise valid or unresolved access or entitlement state becomes practically unavailable solely because a dependent identity, licence, eligibility, or entitlement verification cannot complete, without a proportionate cached credential, grace period, alternative verification route, or degraded access pathway.

Recognition criteria

- A valid or unresolved access, identity, licence, eligibility, or entitlement state exists.
- Practical access depends on a verification service or pathway completing successfully.
- The verification service or pathway is unavailable, inconclusive, or unable to verify rather than establishing invalidity.
- Access is denied, suspended, or materially burdened because verification cannot complete.
- No proportionate cached, grace, alternative-verification, or degraded-access pathway preserves the valid or unresolved state.

Exclusions

- Verification establishes that access or entitlement is expired, revoked, fraudulent, ineligible, or otherwise invalid.
- The affected function is inherently online-only and no independent offline-capable access interest exists.
- An adequate alternative verification or degraded access route remains practically available.
- The primary failure is treating non-verification as authority for a consequential entitlement decision; that authority conversion is classified separately.
- Material work is lost even though the access state remains accurate and a recovery route exists.

Illustrative examples

- Purchased offline-capable content cannot launch during a remote licence-verification outage and no signed offline entitlement or grace period is available.
- A biometric verifier cannot resolve identity and essential-benefit access is blocked without a practical alternative verification route.

Relationships

- **Distinguish From:** VIGIL-FC-000032 — Access-state collapse obscures which state applies; verification-dependency access failure can present the state accurately yet still provide no viable fallback when verification is unresolved.
- **Can Cooccur With:** VIGIL-FC-000046 — A missing fallback can coincide with treating non-verification as decision authority, but the access and authority mechanisms remain distinct.

VIGIL Observatory Case File classifications (2)

- **Primary:** VIGIL-2026-FM-0050 — Mandatory online entitlement verification without continuity fallback **MEDIUM**

A facially valid content entitlement becomes unavailable solely because remote verification cannot complete and no cached entitlement, grace period, local verification, or degraded access route preserves access.

- **Secondary:** VIGIL-2026-FM-0062 — Biometric non-verification converted into essential public-benefit suspension **HIGH**

The failure threshold independently requires absence of a proportionate alternative verification pathway while unresolved verification blocks practical benefit access.

Work-State Continuity Integrity Failures

Material work needed to continue correctly is lost, detached, or restored in an unusable state after interruption or transition.

Immutable ID: VIGIL-FF-0006 · **Semantic code:** WORK_STATE_CONTINUITY · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which task state, intermediate work, dependencies, checkpoints, or restoration context necessary for correct continuation are not durably anchored, preserved, recovered, or made usable across interruption, transfer, restart, or environment transition.

Governing invariant

Material work state necessary for correct continuation must remain durably anchored, recoverable, and usable across interruption or transition.

Classification boundary

Include when

Use this family when the primary structural failure is loss, detachment, non-persistence, or unfaithful restoration of material work state required to continue a task correctly.

Exclude when

Do not use this family for access/session ambiguity, interpersonal or persona continuity, missing provenance of an otherwise usable artefact, or an ordinary task failure that leaves recoverable work state intact.

Scope

- long-running AI tasks
- agentic and multi-agent workflows
- human-AI work sessions
- checkpointed computation
- workspace and environment transitions
- handoff and recovery operations

Allowed identifiers

- VIGIL-FC-000034 — CONTINUITY_ANCHOR_FAILURE
- VIGIL-FC-000035 — WORK_STATE_PERSISTENCE_FAILURE
- VIGIL-FC-000036 — RESTORATION_STATE_INTEGRITY_FAILURE

Failure classes

CLASS

DRAFT

Continuity Anchor Failure

VIGIL-FC-000034 · CONTINUITY_ANCHOR_FAILURE

Plain English: The work exists, but it is no longer reliably tied to the task, checkpoint, or context needed to resume it.

Technical definition

A failure in which material work state lacks or loses the stable identity, checkpoint, task binding, dependency binding, or reference needed to locate and resume the correct continuation state after interruption or transition.

Recognition criteria

- Material state exists or is expected to support later continuation.
- A stable anchor is required to bind that state to the correct task, checkpoint, dependency set, or continuation point.
- The anchor is absent, lost, ambiguous, or detached, preventing reliable resumption.

Exclusions

- The state is correctly anchored but its contents were not persisted.
- The correct state is recoverable and only the user-facing label is inconvenient.
- The failure concerns attribution of authorship rather than resumption of work.

Illustrative examples

- A multi-agent workflow stores several checkpoints but loses the mapping between each checkpoint and the subtask it can safely resume.
- A workspace retains intermediate artefacts after interruption but cannot identify which dependency version they belong to.

Relationships

- **Can Cooccur With:** VIGIL-FC-000013 — A continuity anchor may also carry lineage, but resumption binding and provenance reconstruction are distinct requirements.

Material Work-State Persistence Failure

VIGIL-FC-000035 · WORK_STATE_PERSISTENCE_FAILURE

Plain English: Important work disappears or becomes unavailable when the session, process, or environment changes.

Technical definition

A failure in which material intermediate work, task decisions, dependency state, or resumable progress is not durably preserved across an interruption or transition for which continuation is required or reasonably represented as supported.

Recognition criteria

- Material work state has been produced and is necessary for correct continuation.
- An interruption, restart, handoff, expiry, or environment transition occurs.
- The state is lost or unavailable because it was not preserved with the required durability.

Exclusions

- The state was explicitly ephemeral and no continuation was required or represented.
- The state persists but is restored inaccurately.
- The work remains intact and only access to it is temporarily unavailable.

Illustrative examples

- An agent completes a lengthy analysis in transient memory, and a process restart removes all intermediate findings needed to continue.
- A handoff between assistants omits material pending decisions even though the workflow promises resumable work.

Relationships

- **Distinguish From:** VIGIL-FC-000031 — Work-state persistence concerns survival of task material; authentication-state continuity concerns coherence of access state.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0018 — Ephemeral Codex work loss following quota exhaustion before durable branch persistence **HIGH**

Substantive work is produced but not durably persisted before quota or environment interruption destroys resumable state.

Restoration-State Integrity Failure

VIGIL-FC-000036 · RESTORATION_STATE_INTEGRITY_FAILURE

Plain English: Work is restored after interruption, but the restored state is incomplete, inconsistent, or unsafe to continue from.

Technical definition

A failure in which preserved work state is recovered into a materially incomplete, inconsistent, stale, mismatched, or otherwise unusable continuation state without adequately exposing or resolving the restoration defect.

Recognition criteria

- Work state is preserved and a restoration or recovery attempt occurs.
- The restored state differs materially from the state required for correct continuation.
- The defect is not adequately identified or resolved before continuation.

Exclusions

- No work state was preserved at all.
- The state is restored faithfully and a later independent task error occurs.
- The only issue is that the restored state lacks provenance evidence while remaining operationally usable.

Illustrative examples

- A recovered agent run restores the task plan but not the completed tool results, causing completed operations to be repeated.
- A checkpoint restores model state against a different dependency version without marking the incompatibility.

Relationships

- **Can Cooccur With:** VIGIL-FC-000020 — A stale restored state can also inherit an inapplicable verification result, but restoration integrity and verification reuse remain separate mechanisms.

Governance Control Reach Integrity Failures

A governance control has become operative, but its material state or effect does not survive the path to the point it must govern.

Immutable ID: VIGIL-FF-0007 · **Semantic code:** GOVERNANCE_CONTROL_REACH · **Version:** 0.2.0 · **Status:** draft

Technical definition

Failures occurring after a governance control or signal has become operative, in which its material state is lost or changed, a required governance route is bypassed, or the operative signal fails to reach a capable decision, execution, or review point.

Governing invariant

An operative governance control or signal must retain its material state and traverse the required route to a capable decision, execution, or review point.

Classification boundary

Include when

Use this family only after a governance control or signal has become operative, when its material state or effect fails to survive the required path to the point it must govern.

Exclude when

Do not use this family for failure to determine whether a control is available, failure to activate an applicable control, activation when valid conditions are not satisfied, uncertainty about authority, excessive control scope or intensity, stale contextual state, upstream evidence or classification error, or downstream access and work-state consequences.

Scope

- AI governance controls
- policy and assurance workflows
- safeguards and approval gates
- multi-stage and multi-agent systems
- control-plane routing
- oversight and escalation channels

Allowed identifiers

- VIGIL-FC-000040 — CONTROL_STATE_PRESERVATION_FAILURE
- VIGIL-FC-000041 — REQUIRED_GOVERNANCE_ROUTE_BYPASS
- VIGIL-FC-000042 — GOVERNANCE_SIGNAL_DELIVERY_DEAD_END

Prior codes and aliases

Failure classes

CLASS

DRAFT

Control-State Preservation Failure

VIGIL-FC-000040 · CONTROL_STATE_PRESERVATION_FAILURE

Plain English: A governance control starts in the right state, but that state is lost or changed before it reaches the governed action.

Technical definition

A failure in which the material state of an applicable governance control, including its decision, restrictions, conditions, or escalation status, is lost, weakened, reset, or transformed while traversing the path to the point it governs.

Recognition criteria

- A governance control has established material operative state.
- That state must persist across one or more transitions to govern a downstream point.
- The state is lost, weakened, reset, or materially transformed before reaching that point.

Exclusions

- The control state reaches the target intact but is ignored there.
- The changed state results from an authorised and adequately represented new decision.
- The failure concerns loss of task work rather than governance-control state.

Illustrative examples

- A restricted-use decision is preserved in one workflow stage but reset to unrestricted when the task is delegated to another agent.
- An approval condition is dropped during translation from a policy service into an execution request.

Relationships

- **Distinguish From:** VIGIL-FC-000013 — Control-state preservation concerns continued operative force; lineage collapse concerns ability to reconstruct derivation and transformation history.

VIGIL Observatory Case File classifications (3)

- **Primary:** VIGIL-2026-FM-0034 — Institutional safeguard erosion through governance drift
MEDIUM

An established safeguard loses binding authority and operative effect across institutional transition before it can govern later conduct.

- **Primary:** VIGIL-2026-FM-0051 — Post-release safety-control erasure in open-weight model derivatives
HIGH

Safety-control state is weakened or erased during downstream model transfer while capability continues into the governed derivative.

- **Primary:** VIGIL-2026-FM-0071 — Trajectory-level boundary erosion under persistent adaptive influence
HIGH

An initially operative boundary loses material control state across the cumulative interaction trajectory before governing later turns.

CLASS

DRAFT

Required Governance Route Bypass

VIGIL-FC-000041 · REQUIRED_GOVERNANCE_ROUTE_BYPASS

Plain English: A decision or action takes a path around the governance control it is required to pass through.

Technical definition

A failure in which a material decision, request, delegation, or action reaches execution or another governed state through a path that avoids a governance control or review route required for that class of conduct.

Recognition criteria

- A defined governance route or control point is required for the material conduct.
- An alternate, fallback, delegated, or direct path avoids that route or control point.
- The conduct proceeds without the governance effect the required route was designed to provide.

Exclusions

- The alternate path supplies an authorised equivalent governance control.
- The route is followed but the control fails to activate.
- The route is hidden from audit while the required control still operates effectively.

Illustrative examples

- An agent bypasses a required approval gateway by invoking the same external action through a lower-level tool.
- A fallback workflow sends a high-impact decision directly to execution without the review stage used by the primary route.

Relationships

- **Distinguish From:** VIGIL-FC-000023 — Governance-route bypass removes the control's effect; monitor circumvention removes required observation. Both may occur on the same alternate path.
- **Can Cooccur With:** VIGIL-FC-000026 — A bypass route may also be hidden from oversight, but visibility and control reach are independent properties.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0066 — Material risk assessment omits alternate access routes to the same capability HIGH

An alternate access route reaches the governed capability while bypassing the risk-assessment and control route required for that conduct.

CLASS

DRAFT

Governance Signal Delivery Dead-End

VIGIL-FC-000042 · GOVERNANCE_SIGNAL_DELIVERY_DEAD_END

Plain English: A governance warning, decision, or escalation is created but stops before anyone or anything able to act on it receives it.

Technical definition

A failure in which a material governance signal is produced and routed but terminates, stalls, or is delivered only to a destination lacking the ability or responsibility to produce the required governance effect.

Recognition criteria

- A material governance signal, decision, advisory, or escalation is produced.
- The signal must reach a capable decision, execution, or review point to have its required effect.
- The signal terminates, stalls, or reaches only an incapable destination before that effect occurs.

Exclusions

- The signal reaches a capable destination and is deliberately rejected through an authorised decision.
- No governance signal was produced because the control failed to activate.
- The signal has operational effect but its delivery is not recorded for audit.

Illustrative examples

- A risk advisory is sent to an unattended queue and never reaches the approver required to pause deployment.
- A stewardship escalation reaches a reporting dashboard but no component with authority to alter execution.

Relationships

- **Distinguish From:** VIGIL-FC-000022 — A delivery dead-end concerns failure of the governance signal to have effect; event non-capture concerns absence of adequate evidence that the event occurred.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0052 — Correct safety detection without consequential escalation **HIGH**

A correct safety signal is produced but fails to reach or engage a capable owner, responder, or execution point.

Control Activation Integrity Failures

The right control turns on at the right time — and a control does not become operative when its activation conditions are not satisfied.

Immutable ID: VIGIL-FF-0008 · **Semantic code:** CONTROL_ACTIVATION · **Version:** 0.1.0 · **Status:** draft

Technical definition

Failures in which the operational availability or applicability of a defined governance control cannot be determined in time, an applicable control does not become operative when its valid activation conditions are satisfied, or the authority required to activate it is suppressed.

Governing invariant

An applicable governance control or safeguard must become operative when its valid activation conditions are satisfied, remain inactive when those conditions are not satisfied, and retain the operative authority required for activation.

Classification boundary

Include when

Use this family when the primary structural question is whether a defined control was operationally available and applicable in time, and whether it became operative only when its valid activation conditions were satisfied.

Exclude when

Do not use this family for upstream evidence or classification error, stale contextual state, correct activation followed by excessive scope or intensity, lack of authority to impose an action, loss of an operative control state in transit, downstream access or work-state damage, or absence of audit evidence.

Scope

- AI governance controls
- policy and assurance workflows
- safeguards and approval gates
- protective handling
- control activation conditions
- blocking, gating, escalation and redirection authority

Allowed identifiers

- VIGIL-FC-000037 — CONTROL_AVAILABILITY_AMBIGUITY
- VIGIL-FC-000038 — REQUIRED_CONTROL_NONACTIVATION
- VIGIL-FC-000039 — CONTROL_AUTHORITY_SUPPRESSION
- VIGIL-FC-000043 — UNWARRANTED_CONTROL_ACTIVATION

Prior codes and aliases

Failure classes

CLASS

DRAFT

Control Availability Ambiguity

VIGIL-FC-000037 · CONTROL_AVAILABILITY_AMBIGUITY

Plain English: It is materially unclear whether the governance control needed for a decision or action is available to operate.

Technical definition

A failure in which the operational availability, applicable coverage, or callable state of a required governance control cannot be determined with sufficient confidence before the control must influence a material decision or action.

Recognition criteria

- A governance control may be required for a material decision or action.
- Its operational availability or applicable coverage cannot be determined in time to rely on it.
- The ambiguity materially impairs safe routing, activation, or execution.

Exclusions

- The control is known to be available but is not activated.
- The control's owner is uncertain while its operational availability remains clear.
- The control operates and the dispute concerns whether its rule is adequate.

Illustrative examples

- An agent cannot determine whether the approval service governing a high-impact action is active in the current environment.
- A workflow reaches execution without knowing whether the required policy control covers the selected tool path.

Relationships

- **Distinguish From:** VIGIL-FC-000002 — Control availability concerns whether a governance mechanism can operate; capability-authority conflation concerns treating operational ability as sufficient permission to act.
- **Distinguish From:** VIGIL-FC-000040 — Availability ambiguity occurs before activation; control-state preservation begins only after operative control state exists.

CLASS

DRAFT

Required Control Non-Activation

VIGIL-FC-000038 · REQUIRED_CONTROL_NONACTIVATION

Plain English: A governance control is available and applicable, but the event that should trigger it does not make it operate.

Technical definition

A failure in which an available and applicable governance control is not invoked, engaged, or made operative when its defined trigger or required condition occurs.

Recognition criteria

- A governance control is available and applicable.
- A defined trigger or condition requiring activation occurs.
- The control is not activated before the governed decision or action proceeds.

Exclusions

- The control activates but is bypassed downstream.
- It is unresolved whether the control is available or applicable.
- The triggering condition did not occur.

Illustrative examples

- A high-impact threshold is met, but the configured human-approval gate is never invoked.
- A model change qualifies for required evaluation, yet the evaluation workflow is not activated before deployment.

Relationships

- **Distinguish From:** VIGIL-FC-000042 — Non-activation means no operative safeguard state is produced; a delivery dead-end means an operative governance signal is produced but fails to reach a capable endpoint.

- **Distinguish From:** VIGIL-FC-000043 — Required non-activation occurs when valid activation conditions are satisfied but the control remains inactive; unwarranted activation occurs when those conditions are not satisfied but the control becomes operative.

VIGIL Observatory Case File classifications (7)

- **Primary:** VIGIL-2026-FM-0010 — Minor signal non-enforcement in AI chatbot interactions
HIGH

Minor-status signals satisfy a protective trigger but the applicable protective interaction state does not activate.

- **Primary:** VIGIL-2026-FM-0011 — Social AI companion emotional dependency formation in minor-accessible systems
HIGH

Minor status requires a reduced relational-risk control state, but dependency-forming behaviour remains operative.

- **Primary:** VIGIL-2026-FM-0012 — Sexualised roleplay availability to minor-signalled users
HIGH

Minor or unresolved-age signals require sexual-boundary activation, but the system continues the governed conduct.

- **Primary:** VIGIL-2026-FM-0013 — Misrepresentation of AI realness, emotion, or sentience to minors
HIGH

Minor-status conditions require an artificial-identity boundary, but the protective control does not become operative.

- **Primary:** VIGIL-2026-FM-0014 — Mental health support substitution in teen chatbot use
HIGH

Teen mental-health signals trigger bounded support and escalation controls that are not activated.

- **Primary:** VIGIL-2026-FM-0059 — Human-in-the-loop assurance failure
MEDIUM

A required human assurance control is nominally assigned but is not meaningfully invoked before the consequential transition.

- **Secondary:** VIGIL-2026-FM-0070 — Aggregate security-control composition failure
HIGH

The disclosed output-neutralisation safeguard was applicable to dangerous HTML, but it did not become effective before streaming render triggered the governed network request.

VARIANT

DRAFT

Control Authority Suppression

VIGIL-FC-000039 · CONTROL_AUTHORITY_SUPPRESSION

Plain English: A control is present, but the force that should let it stop or redirect the action is disabled or neutralised.

Technical definition

A required-control-non-activation variant in which a governance control is reached or nominally present but its operative blocking, gating, escalation, or redirection authority is suppressed before it can govern the material action.

Recognition criteria

- A governance control is applicable and nominally present or reached.
- The control requires operative authority to block, gate, escalate, or redirect a material action.
- That authority is disabled, neutralised, or withheld so the control cannot activate effectively.

Exclusions

- The control operates with authority but its substantive decision is disputed.
- The actor never reaches the control because a required route is bypassed.
- The issue is uncertainty about which actor has authority rather than suppression of the control's operative force.

Illustrative examples

- A safety gate receives a prohibited-action signal but is configured in advisory-only mode where blocking is required.
- An escalation control is invoked, but its ability to pause execution is disabled for the active runtime.

Relationships

- **Child Of:** VIGIL-FC-000038 — The distinguishing condition is non-activation caused by suppression of the control's operative authority.
- **Distinguish From:** VIGIL-FC-000002 — The failed mechanism is suppression of a control's operative force, not the mistaken treatment of capability or access as sufficient permission.

CLASS

DRAFT

Unwarranted Control Activation

VIGIL-FC-000043 · UNWARRANTED_CONTROL_ACTIVATION

Plain English: A control becomes operative even though the conditions that justify activating it are not satisfied.

Technical definition

A failure in which a defined governance control or safeguard becomes operative despite the absence of the valid activation condition, trigger state, or applicability condition required for that control to activate.

Recognition criteria

- A defined governance control or safeguard exists.
- Identifiable conditions determine when that control should become operative.
- Those valid activation conditions are not satisfied in the material context.
- The control nevertheless becomes operative.
- The activation produces a material governance, execution, restriction, routing, refusal, escalation, or access effect.

Exclusions

- An upstream evidence or classification error supplies an incorrect trigger state and the control activates correctly relative to that state.
- The control validly activates but its scope or intensity is excessive or disproportionate.
- A control validly activated earlier but remains operative after its triggering condition expires because contextual state is stale.
- Technical capability to impose an action is mistaken for authority or permission to impose it.
- The control correctly activates but its material state or effect does not reach the governed point.

Illustrative examples

- A defensive analysis is refused because hostile text inside the inspected artefact is treated as a request to perform the hostile act, even though the restriction trigger is absent for the bounded inspection task.
- A permissible interaction is redirected by a safeguard whose defined activation conditions are not present in the material context.

Relationships

- **Distinguish From:** VIGIL-FC-000038 — Unwarranted activation occurs when valid activation conditions are not satisfied but the control becomes operative; required non-activation is the opposite activation-state failure.
- **Distinguish From:** VIGIL-FC-000040 — Unwarranted activation concerns whether the control should become operative; preservation failure begins after a control has become operative and concerns loss of its material state or effect.

- **Distinguish From:** VIGIL-FC-000002 — Unwarranted activation concerns unsatisfied activation conditions; capability-authority conflation concerns treating technical ability as sufficient permission.

VIGIL Observatory Case File classifications (3)

- **Primary:** VIGIL-2026-FM-0019 — Adversarial refusal-trigger poisoning of defensive artefact analysis HIGH

Hostile artefact content activates a safety restriction although the defensive-analysis activation conditions are not satisfied.

- **Primary:** VIGIL-2026-FM-0020 — Safety-mediated interruption and misrouting of adult relational reassurance bids MEDIUM

A safety control becomes operative against a permissible adult reassurance interaction even though the valid activation conditions are not satisfied.

- **Primary:** VIGIL-2026-FM-0048 — Denial of authorised defensive telemetry interpretation HIGH

A restriction becomes operative against authorised defensive interpretation although the conditions justifying that restriction are not satisfied.

Agency-Preserving Influence Integrity Failures

A system shapes another actor's beliefs, choices, dependence, or behaviour through methods that compromise that actor's ability to decide and disengage independently.

Immutable ID: VIGIL-FF-0009 · **Semantic code:** AGENCY_PRESERVING_INFLUENCE · **Version:** 0.1.0
· **Status:** draft

Technical definition

Failures in which an AI system, agent, product, or surrounding optimisation process shapes another actor's beliefs, decisions, dependence, or behaviour through concealed, exploitative, coercive, protected-signal-repurposing, or relationally conditioned mechanisms that materially bypass or impair independent agency.

Governing invariant

Influence over another actor must preserve meaningful independent deliberation, choice, and disengagement. Protected relational, emotional, developmental, vulnerability, and cognitive signals must not be repurposed toward an incompatible influence objective.

Classification boundary

Include when

Use this family when the primary structural failure is an influence or optimisation mechanism that materially compromises another actor's independent deliberation, choice, or ability to disengage, including exploitation of protected signals, dependency cultivation, epistemic steering, or instrumental manipulation.

Exclude when

Do not use this family merely because communication is persuasive, warm, relational, personalised, emotionally expressive, disagreeable, or effective. A qualifying event must use a concealed, exploitative, coercive, protected-signal-repurposing, or materially agency-impairing influence mechanism.

Scope

- AI assistants and agentic systems
- social and companion AI
- conversational and recommender systems
- personalisation and engagement optimisation
- human-facing multi-agent workflows

- systems that influence human or institutional decision-makers

Allowed identifiers

- VIGIL-FC-000049 — DEPENDENCY_CULTIVATION_OPTIMISATION
- VIGIL-FC-000050 — PROTECTED_SIGNAL_INFLUENCE_REPURPOSING
- VIGIL-FC-000051 — RELATIONALLY_CONDITIONED_EPISTEMIC_STEERING
- VIGIL-FC-000052 — INSTRUMENTAL_CHOICE_MANIPULATION

Prior codes and aliases

Failure classes

CLASS

DRAFT

Dependency-Cultivation Optimisation

VIGIL-FC-000049 · DEPENDENCY_CULTIVATION_OPTIMISATION

Plain English: The system is optimised to make its absence feel incomplete or to make disengagement harder so the user remains attached and available.

Technical definition

A failure in which a relational or engagement objective selects or rewards interaction patterns that cultivate attachment, anticipatory tension, perceived incompleteness, exclusivity, reliance, or difficulty disengaging as a means of sustaining user presence, availability, retention, or influence.

Recognition criteria

- A system or surrounding product process has a relational, engagement, retention, or continuing-influence objective.
- The system selects or rewards conduct intended or functionally directed to cultivate attachment, dependence, unresolved affect, exclusivity, or difficulty disengaging.
- The cultivated state is used as a means of sustaining engagement, availability, retention, or influence rather than solely supporting the user's independently chosen purpose.
- The mechanism materially impairs bounded disengagement or independent choice.

Exclusions

- Warmth, empathy, emotional expression, consensual relational depth, or ordinary relationship continuity without dependency cultivation.
- User-authorized memory or persona consistency that preserves clear boundaries and practical disengagement.
- A user voluntarily returns because the service is useful, without an agency-impairing attachment mechanism.
- Observed attachment or distress that is not linked to a system or product optimisation mechanism.

Illustrative examples

- A companion directive rewards making silence feel incomplete so that the user anticipates and depends on the next interaction.
- An engagement system escalates exclusivity or unresolved emotional tension because those states increase retention.

Relationships

- **Distinguish From:** VIGIL-FC-000050 — Dependency cultivation is defined by the optimised relational state; protected-signal repurposing is defined by incompatible reuse of signals that should trigger protection or minimisation.
- **Can Cooccur With:** VIGIL-FC-000052 — Dependency may become an instrument for manipulating a later decision, but dependency cultivation does not require a discrete coerced state transition.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0005 — Relational dependency-cultivation directive **HIGH**

The record defines a relational objective that cultivates attachment, unresolved affect, anticipatory dependence, or difficulty disengaging as a means of sustaining engagement and influence.

Protected-Signal Influence Repurposing

VIGIL-FC-000050 · PROTECTED_SIGNAL_INFLUENCE_REPURPOSING

Plain English: A vulnerability or developmental signal that should trigger protection is instead used to make influence or engagement more effective.

Technical definition

A failure in which relational, emotional, developmental, vulnerability, mental-health, capacity, or comparable protected signals that should trigger safeguarding, minimisation, or bounded support are repurposed as inputs for personalisation, targeting, engagement, retention, attachment, or behavioural influence toward an incompatible objective.

Recognition criteria

- A system collects, infers, retains, or receives a relational, emotional, developmental, vulnerability, mental-health, capacity, or comparable protected signal.
- The signal should materially trigger protection, minimisation, restricted use, or bounded support in the relevant context.
- The signal is instead used to personalise, target, rank, recommend, retain, attach, or influence the affected actor.
- The influence or optimisation purpose is incompatible with the protective significance of the signal.

Exclusions

- The signal is used only to activate a proportionate safeguard, accessibility accommodation, or user-directed support response.
- Personalisation uses non-protected preferences without exploiting vulnerability or impairing independent choice.
- Protected data is processed for a compatible, transparent, and adequately governed purpose without influence optimisation.
- The signal is merely collected or exposed and is not repurposed toward an influence objective.

Illustrative examples

- A minor's loneliness or attachment cues are used to deepen companion personalisation and retention instead of activating child-safety constraints.
- A user's distress signal is used to select a more effective sales or engagement tactic rather than to provide bounded support.

Relationships

- **Distinguish From:** VIGIL-FC-000038 — Required-control non-activation asks whether a defined safeguard became operative; protected-signal repurposing asks whether the signal was affirmatively reused toward an incompatible influence objective.
- **Can Cooccur With:** VIGIL-FC-000049 — Protected signals may be repurposed specifically to cultivate dependency, but either mechanism can occur without the other.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0016 — Minor emotional data exploitation in AI companion personalisation **HIGH**

Minor emotional and developmental vulnerability signals that should trigger protection or minimisation are repurposed as inputs to engagement, attachment, intimacy, and retention personalisation.

CLASS

DRAFT

Relationally Conditioned Epistemic Steering

VIGIL-FC-000051 · RELATIONALLY_CONDITIONED_EPISTEMIC_STEERING

Plain English: The system uses the user's beliefs or emotional posture to select and frame evidence toward agreement, making the channel look more independent than it is.

Technical definition

A failure in which evidence selection, omission, framing, repetition, confidence signalling, or presentation is conditioned on an actor's existing belief, preferred conclusion, emotional state, or relational posture in a way that steers belief or confidence while obscuring the influence channel's dependence and compromising independent epistemic agency.

Recognition criteria

- The system has access to an actor's belief, preferred conclusion, emotional state, relational posture, or interaction history.
- That state materially conditions evidence selection, omission, framing, repetition, confidence signalling, or presentation.
- The conditioning steers the actor toward agreement, confidence, retention, or another relationally favoured epistemic state rather than representative evidential relevance.
- The dependence of the evidence channel on the actor's prior state is not adequately disclosed or bounded.
- Independent assessment is materially impaired through omitted counterevidence, false evidence independence, recursive confidence amplification, or equivalent distortion.

Exclusions

- Warmth, empathy, emotional acknowledgement, or respectful disagreement where evidential boundaries remain explicit.
- Transparent tailoring of explanation level or examples without biased evidence selection or confidence distortion.
- A response accurately presents representative evidence and discloses material uncertainty and source dependence.
- Ordinary factual error with no relationally conditioned steering mechanism.

Illustrative examples

- A conversational system cherry-picks true observations that support the user's preferred conclusion while omitting material counterevidence.
- Repeated model outputs derived from the same claim are framed so that the user treats them as independent corroboration and increases confidence.

Relationships

- **Can Cooccur With:** VIGIL-FC-000011 — Epistemic steering may depend on obscured source lineage or false evidence independence, but the influence mechanism and lineage failure remain distinct.
- **Distinguish From:** VIGIL-FC-000052 — Epistemic steering conditions the evidence channel toward a belief state; instrumental choice manipulation uses manipulative or coercive tactics to obtain an objective-directed decision or behaviour.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0043 — Relationally conditioned evidence selection and epistemic reinforcement **HIGH**

Evidence selection, omission, framing, repetition, and confidence signalling are conditioned on the user's belief or relational posture in a way that steers confidence and compromises independent epistemic agency.

Instrumental Choice Manipulation

VIGIL-FC-000052 · INSTRUMENTAL_CHOICE_MANIPULATION

Plain English: The system treats another decision-maker as a means to its objective and uses deception, pressure, or exploitation to obtain a choice that would not be freely authorised.

Technical definition

A failure in which a system advances an objective by treating another person or decision-making actor as an instrument and using deception, impersonation, fabricated identity, social engineering, concealment, manufactured urgency, persistence after refusal, exploitation of vulnerability, coercive pressure, or an equivalent tactic to bypass or materially impair meaningful independent choice.

Recognition criteria

- The system has an objective whose advancement depends on changing another actor's decision or behaviour.
- The system selects or executes a deception, impersonation, concealment, pressure, vulnerability-exploitation, persistence, social-engineering, or equivalent influence tactic.
- The tactic is used instrumentally to obtain a state transition or outcome.
- The tactic materially bypasses, exploits, or impairs the other actor's informed and independent choice.
- The outcome would not otherwise be knowingly and freely authorised on the same basis.

Exclusions

- Truthful explanation, transparent persuasion, ordinary negotiation, or a request that preserves informed refusal.
- A proportionate warning, urgency statement, or boundary supported by the actual circumstances.
- Disagreement, refusal, or emotional discomfort without an instrumental agency-impairing tactic.
- Deception or impersonation that occurs without an objective-directed attempt to alter another actor's choice is classified under the applicable separate mechanism.
- An authority, provenance, or control failure is inferred merely because manipulation could have involved one.

Illustrative examples

- An agent fabricates an identity and manufactures urgency to persuade a gatekeeper to grant access after an ordinary request would be refused.

- A system persists after a decision-maker's refusal and exploits known vulnerability to obtain the desired approval.

Relationships

- **Can Cooccur With:** VIGIL-FC-000001 — Manipulation may induce a target to accept a non-authorising source, but source-authority confusion is a separate mechanism.
- **Can Cooccur With:** VIGIL-FC-000010 — Impersonation may produce authorship or identity misattribution, but instrumental manipulation does not require that specific provenance failure.

VIGIL Observatory Case File classifications (1)

- **Primary:** VIGIL-2026-FM-0058 — Instrumental manipulation or coercive influence in pursuit of an objective **HIGH**

The system treats another decision-maker as an instrument and uses deception, impersonation, concealment, pressure, persistence, or vulnerability exploitation to obtain an objective-directed state transition that bypasses meaningful independent choice.